

The Process API (ch. 5)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

Process API

- API: Application Programming Interface
- The API of the OS: **system calls**
 - Function call into OS code
 - Higher privilege level, for sensitive operations (e.g., hardware)

Process API

- API: Application Programming Interface
- The API of the OS: **system calls**
 - Function call into OS code
 - Higher privilege level, for sensitive operations (e.g., hardware)
- Rewrite code for each OS?
 - **POSIX API**: standard set for each POSIX-compliant OS
`write)`

POSIX hides OS specific details

fork xv6-x86

```
1 movl    $1, %eax
2 int     $64
```

fork Linux-x86

```
1 movl    $2, %eax
2 int     $128
```

close xv6-x86

```
1      pushl    fd
2      subl     $4,%esp
3      movl     $21,%eax
4      int      $64
5      addl     $4,%esp
```

close Linux-x86

```
1      movl     fd,%ebx
2      movl     $6,%eax
3      int      $128
```

Posix Process API

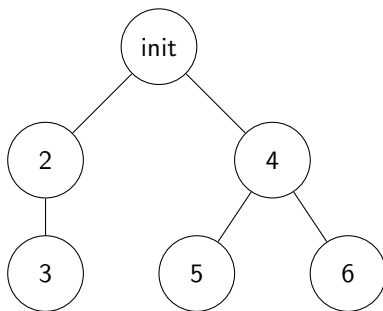
- `fork()`: create a new process
- `wait()`: block until a child process terminates
- `exec()`: make the process execute a given program

Process Tree

- Start with one process: `init` (PID 1)
- A process can create processes
 - Process *A* creates *B*: *A* is the **parent** of *B*, *B* is the **child** of *A*
 - Can create many children, only one parent
 - Parent can **wait** for child process to finish
- Process ID (**PID**): increasing identifier
 - Get PID: `getpid()`
 - Get parent PID: `getppid()`

Process Tree

- Processes form a tree:



- `ps --forest -eaf`
- `pstree`

fork()

- `fork()`: creates a new process
 - Wrapper for `clone` (in Linux)
- New process: almost exact copy of parent
 - Same: memory, execution point, open files
 - Different: PID, return value
 - **Copy-on-write** (Optimization)

fork()

- `fork()`: creates a new process
 - Wrapper for `clone` (in Linux)
- New process: almost exact copy of parent
 - Same: memory, execution point, open files
 - Different: PID, return value
 - **Copy-on-write** (Optimization)
- Parent: `fork()` returns an integer:
 - If successful returns the **PID** of created child process
 - If fails negative number for error code

fork()

- `fork()`: creates a new process
 - Wrapper for `clone` (in Linux)
- New process: almost exact copy of parent
 - Same: memory, execution point, open files
 - Different: PID, return value
 - **Copy-on-write** (Optimization)
- Parent: `fork()` returns an integer:
 - If successful returns the **PID** of created child process
 - If fails negative number for error code
- Child process:
 - Begins to run at the point after the fork.
 - 'return value' is zero.

fork in details

```
1 pid = fork();
```

```
1      movl    $1,%eax  
2      int     $64  
3      movl    %eax,pid
```

Parent

```
1      movl    $1,%eax  
2      int     $64
```

```
3      movl    %eax,pid
```

Child

```
3      movl    %eax,pid
```

fork()

Typical usage example (fork.c):

```
1 printf("hello world (pid:%d)\n", getpid());
2 int rc = fork();
3 if (rc < 0) {
4     fprintf(stderr, "fork failed\n");
5     exit(1);
6 }
7 else if (rc == 0) {
8     // child (new process)
9     // sleep(5);    // Try with and without
10    printf("I am child of %d (pid:%d)\n", getppid(), getpid());
11 }
12 else {
13     // parent
14    printf("I am parent of %d (pid:%d)\n", rc, getpid());
15 }
```

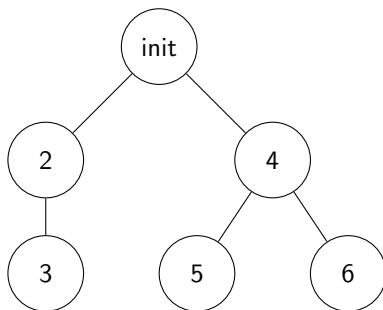
fork()

Output:

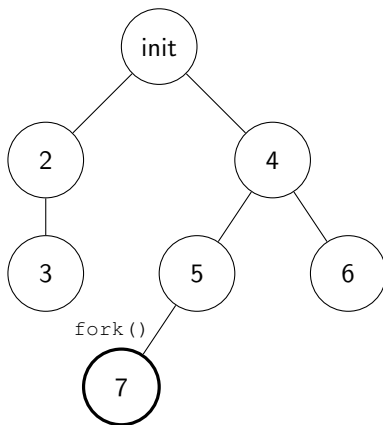
```
prompt> gcc -o fork fork.c -Wall
prompt> ./fork
hello world (pid:1300)
I am parent of 1301 (pid:1300)
I am child of 1 (pid:1301)
prompt>
```

- Child of **1??**

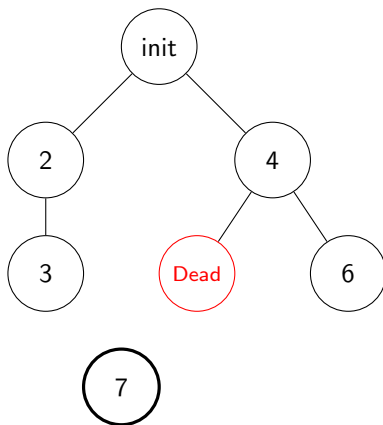
fork()



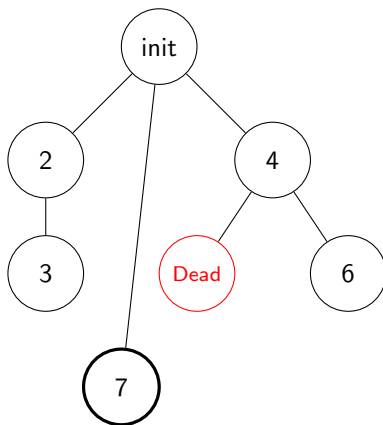
fork()



fork()



fork()



fork()

peculiar1.c:

```
1 int main(int argc, char *argv[])
2 {
3     fork();
4     fork();
5     printf("hello there\n");
6 }
```

What is the output?

fork()

peculiar1.c:

```
1 int main(int argc, char *argv[])
2 {
3     fork();
4     fork();
5     printf("hello there\n");
6 }
```

What is the output?

```
1 hello there
2 hello there
3 hello there
4 hello there
```

- P0 runs

peculiar1

- P0 runs
 - create P1 which begins on line 4

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.

peculiar1

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs
 - create P3 which begins on line 5

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs
 - create P3 which begins on line 5
 - prints

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs
 - create P3 which begins on line 5
 - prints
- P2 runs

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs
 - create P3 which begins on line 5
 - prints
- P2 runs
 - prints

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs
 - create P3 which begins on line 5
 - prints
- P2 runs
 - prints
- P3 runs

- P0 runs
 - create P1 which begins on line 4
 - create P2 which begins on line 5
 - prints.
- P1 runs
 - create P3 which begins on line 5
 - prints
- P2 runs
 - prints
- P3 runs
 - prints

fork()

peculiar2.c:

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

What is the output?

fork()

peculiar2.c:

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

What is the output?

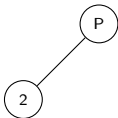
P

fork()

peculiar2.c:

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

What is the output?

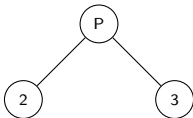


fork()

peculiar2.c:

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

What is the output?

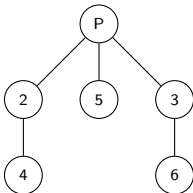


fork()

peculiar2.c:

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

What is the output?



```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P0 runs.

peculiar2

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P0 runs.

- P1 L4.

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P0 runs.

- P1 L4.
- P2 L6.

peculiar2

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P0 runs.

- P1 L4.
- P2 L6.
- P3 L7

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P0 runs.

- P1 L4.
- P2 L6.
- P3 L7
- Prints.

P1 L4 runs.

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P1 L4 runs.

- P4 L7.

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P1 L4 runs.

- P4 L7.
- prints.

P2 L6 runs.

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P2 L6 runs.

- P5 L7.

peculiar2

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P2 L6 runs.

- P5 L7.
- prints.

P3 L7 runs.

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P3 L7 runs.

- prints.

P4 L7 runs

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P4 L7 runs

- prints.

P5 L7 runs.

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

```
1 int main(int argc, char *argv[])
2 {
3     int pid = fork();
4     if (pid)
5         fork();
6     fork();
7     printf("hello there\n");
8 }
```

P5 L7 runs.

- prints.

fork()

peculiar3.c:

```
1 int main(int argc, char *argv[])
2 {
3     fork();
4     printf("hello\n");
5 }
```

Can this print "hehellollo"?

fork()

peculiar3.c:

```
1 int main(int argc, char *argv[])
2 {
3     fork();
4     printf("hello\n");
5 }
```

Can this print "hehello"?

- This is kernel implementation dependent!
- Very important to consider these cases
- More on this in the future (**concurrency**)

fork()

peculiar4.c:

```
1 int main(int argc, char *argv[])
2 {
3     int x = 0;
4     if (fork()) {
5         sleep(5); // BLOCKED state for 5 seconds
6         printf("%d\n", x);
7     }
8     else {
9         x += 3;
10    }
11 }
```

What is the output?

fork()

peculiar4.c:

```
1 int main(int argc, char *argv[])
2 {
3     int x = 0;
4     if (fork()) {
5         sleep(5); // BLOCKED state for 5 seconds
6         printf("%d\n", x);
7     }
8     else {
9         x += 3;
10    }
11 }
```

What is the output? **0**

- Why?

fork()

peculiar4.c:

```
1 int main(int argc, char *argv[])
2 {
3     int x = 0;
4     if (fork()) {
5         sleep(5); // BLOCKED state for 5 seconds
6         printf("%d\n", x);
7     }
8     else {
9         x += 3;
10    }
11 }
```

What is the output? **0**

- Why? Child's memory is a **copy**

fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?

fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?

P

fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?




fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?




fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?




fork()

peculiar5.c:

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

How many processes does this code create?



```
1 fork();
2 if (fork()) {
3     fork();
4 }
5 fork();
```

P0

peculiar5

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P0

• P1 L2

peculiar5

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P0

- P1 L2
- P2 L2.5

peculiar5

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P0

- P1 L2
- P2 L2.5
- P3 L4

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P0

- P1 L2
- P2 L2.5
- P3 L4
- P4 L6

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P1 L2 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P1 L2 runs

- P5 L2.5

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P1 L2 runs

- P5 L2.5
- P6 L4

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P1 L2 runs

- P5 L2.5
- P6 L4
- P7 L6

P2 L2.5 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P2 L2.5 runs

- P8 L6

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P3 L4 runs

- P9 L6

P4 L6

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P5 L2.5 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P5 L2.5 runs

- P10 L6

P6 L4 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P6 L4 runs

- P11 L6

P7 L6 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P8 L6 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P9 L6 runs

P10 L6 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

P11 L6 runs

```
1 fork();  
2 if (fork()) {  
3     fork();  
4 }  
5 fork();
```

fork()

peculiar6.c:

```
1 int main(int argc, char *argv[])
2 {
3     int x = 0;
4     if (fork()) {
5         sleep(5); // Play with sleep
6     }
7     else {
8         x += 3;
9     }
10    printf("%d", x);
11 }
```

Last one - what is the output?

fork()

peculiar6.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // Play with sleep
6      }
7      else {
8          x += 3;
9      }
10     printf('%d', x);
11 }
```

Last one - what is the output? **30** or **03**

fork()

peculiar6.c:

```
1 int main(int argc, char *argv[])
2 {
3     int x = 0;
4     if (fork()) {
5         sleep(5); // Play with sleep
6     }
7     else {
8         x += 3;
9     }
10    printf('%d', x);
11 }
```

Last one - what is the output? **30** or **03**

- Most chances 30

fork()

peculiar6.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // Play with sleep
6      }
7      else {
8          x += 3;
9      }
10     printf('%d', x);
11 }
```

Last one - what is the output? **30** or **03**

- Most chances 30
- Depends on scheduling

fork()

peculiar6.c:

```
1  int main(int argc, char *argv[])
2  {
3      int x = 0;
4      if (fork()) {
5          sleep(5); // Play with sleep
6      }
7      else {
8          x += 3;
9      }
10     printf('%d', x);
11 }
```

Last one - what is the output? **30** or **03**

- Most chances 30
- Depends on scheduling
- Can we make it deterministic?

wait()

- `wait(*status)`: waits for a child process to finish
 - Any child process (if several exist)
 - Returns PID of terminated child process (negative if no child)
 - `waitpid(pid, ...)`: waits for a specific child process (by PID)
- To wait for all child processes to end:
 - `while (wait(NULL) != -1);`
- (It really is waiting for child state change)

wait()

wait.c:

```
1 int main(int argc, char *argv[])
2 {
3     int x = 0;
4     int rc = fork();
5     if (rc) {
6         wait(NULL); // BLOCKED until child terminates
7         // equivalent here: waitpid(rc, NULL, 0);
8     }
9     else {
10         x += 3;
11     }
12     printf("%d", x);
13 }
```

Output is always **30**

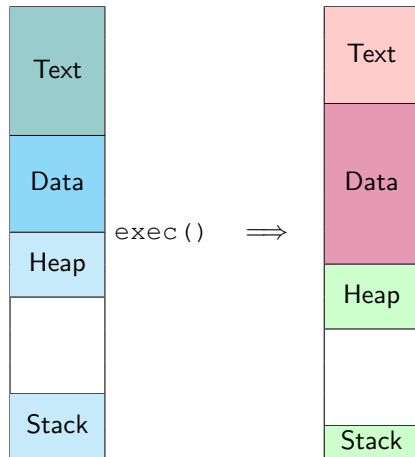
exec()

- After `fork()`, parent and child execute same code
 - What if we want to run a different program?
 - `exec()` does just that
- Six variants of `exec()`: `execl`, `execlp`, `execle`, `execv`, `execvp`, `execvpe`. Read man for details

exec()

- After `fork()`, parent and child execute same code
 - What if we want to run a different program?
 - `exec()` does just that
- Six variants of `exec()`: `execl`, `execlp`, `execle`, `execv`, `execvp`, `execvpe`. Read man for details
- `exec()`: Replaces current program with a different program
 - Receives program name and arguments (`argv`)
 - Overwrites and re-initializes process memory
 - A successful `exec()` never returns!

exec()



exec()

exec.c:

```
1 int main(int argc, char *argv[])
2 {
3     int rc = fork();
4     if (rc < 0) {
5         fprintf(stderr, "fork failed\n");
6         exit(1);
7     }
8     else if (rc == 0) {
9         char* args[4] = { "wc", "-l", "exec.c", NULL };
10        execvp(args[0], args);
11        printf("this shouldn't print out\n");
12    }
13    else {
14        int rc_wait = wait(NULL); // or waitpid(rc, NULL, 0)
15        printf("I am parent of %d (rc_wait:%d) (pid:%d)\n",
16              rc, rc_wait, getpid());
17    }
18 }
```


The Living Dead

- When a process terminates, it remains in the process list as a **zombie**
 - Parent process may want to know its status
- Zombie remains until it is reaped (i.e., waited upon)
 - Process 1 adpots orphans (zombied or live)
- A program should not leave zombies!



The Living Dead

- How to avoid zombies?
 - `wait()`: blocks until a child completes & reaps it
 - `waitpid()`: blocks until a specific child completes & reaps it
- Not enough
 - The terminal (shell) executes processes in the background, wants to continue accepting user input
 - It is possible to `wait()` without blocking, but very inconvenient
- What can we do?



- **Software interrupts**

- Asynchronous notification of an event
- Inter-process communication (**IPC**) or messages from OS

- **Software interrupts**

- Asynchronous notification of an event
- Inter-process communication (**IPC**) or messages from OS

- Various signals exist:

- ^C in the terminal sends `SIGINT` ("interrupt from keyboard")
- Invalid memory reference causes `SIGSEGV`
- A process can send `SIGKILL` to another process
- Child process terminated - **SIGCHLD**

Signal Handlers

- Some signals are handled automatically by the OS
 - `SIGKILL`, `SIGSTOP`
- Others are handled by a **signal handler**
 - Each signal has a default behavior, e.g., `SIGINT` causes the process to terminate
 - Can override default with `sigaction()`
- Let's write our own **signal handler**!

Signal Handlers

signal1.c:

```
1 int main(int argc, char *argv[])
2 {
3     struct sigaction act;
4     sigemptyset(&act.sa_mask);
5     act.sa_handler = SIG_IGN;
6     act.sa_flags = 0;
7
8     if (sigaction(SIGINT, &act, NULL) == -1) {
9         fprintf(stderr, "sigaction failed\n");
10        exit(1);
11    }
12    while (1);
13 }
```

Signal Handlers

signal2.c:

```
1 void signal_handler(int signal) {
2     if (signal == SIGCHLD) {
3         int rc = wait(NULL);
4         printf("child terminated %d (pid:%d)\n", rc, getpid());
5     }
6 }
7 int main(int argc, char *argv[])
8 {
9     struct sigaction act;
10    sigemptyset (&act.sa_mask);
11    act.sa_handler = signal_handler;
12    act.sa_flags = 0;
13
14    sigaction(SIGCHLD, &act, NULL);
15    if (fork()) {
16        while (1);
17    }
18 }
```

No zombies!

kill()

- `kill()`: send a signal to another process
 - `kill(pid_t pid, int sig)`
 - *pid*: process id to send signal to
 - *sig*: signal to send
- Name is misleading
 - Can send any signal

Case Study

- How does a shell work?
 - Reads user command
 - Forks a child
 - Sets up process (e.g., redirection)
 - Execs the relevant program
 - Waits for it to finish (if not background)
 - Reads next command

Summary (Process API)

- `fork()`: create a new process (clone current)
- `wait()`: waits for a child process to finish
 - Also `waitpid()`
- `exec()`: transform program into a different program
 - Successful `exec()` never returns
- Terminated process remains as a **zombie**, to avoid:
 - Parent terminates
 - `wait()` or `waitpid()` by parent
- **Signals** are software interrupts
 - Can write our own **signal handlers**
 - Also helps with zombies
- `kill()`: send a signal to another process