

# Limited Direct Execution (ch. 6)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

# Virtualizing the CPU

- Virtualizing the CPU by **time sharing**
- Two challenges:
  - **Performance:** avoid adding excessive overhead
  - **Control:** retain control over the CPU
- Without control, a process could:
  - Run forever and take over the machine
  - Access information it should not be allowed to access

# Direct Execution

- **Direct execution:** run the program directly on the CPU

# Direct Execution

- **Direct execution:** run the program directly on the CPU

OS	Program
Create entry for process list Allocate memory for program Load program into memory Set up stack with <code>argc/argv</code> Clear registers Execute <b>call</b> <code>main()</code>  Free memory of process Remove from process list	         Run <code>main()</code> Execute <code>return from main()</code>

# Direct Execution

This approach is problematic:

- ❶ Full access to hardware
  - A bug could corrupt the hardware, freeze the machine, etc.
- ❷ How does the OS stop a running process?
  - Implementing **time sharing**
  - What if a process goes into an infinite loop?

# Direct Execution

This approach is problematic:

- ❶ Full access to hardware
  - A bug could corrupt the hardware, freeze the machine, etc.
- ❷ How does the OS stop a running process?
  - Implementing **time sharing**
  - What if a process goes into an infinite loop?

Without limits on processes, the OS is not in control of anything

# Direct Execution

This approach is problematic:

- 1 Full access to hardware
  - A bug could corrupt the hardware, freeze the machine, etc.
- 2 How does the OS stop a running process?
  - Implementing **time sharing**
  - What if a process goes into an infinite loop?

Without limits on processes, the OS is not in control of anything

The solution: **limited** direct execution

# Restricted Operations

- New processor mode: **user mode**
  - Code in user mode is restricted
  - e.g., no I/O requests
- OS runs in **kernel mode**
  - Code in kernel mode can do what it likes



# Restricted Operations

- New processor mode: **user mode**
  - Code in user mode is restricted
  - e.g., no I/O requests
- OS runs in **kernel mode**
  - Code in kernel mode can do what it likes
- What if a process wishes to perform a privileged operation?

# System Calls

- Carefully expose functionality to user programs
- API for sensitive kernel operations
- Each system call has a known number

# System Calls

- Classically, special **trap** instruction (Some modern processors offer less heavy methods)
  - Raises privilege level to **kernel mode**
  - Jumps into the kernel code
  - The **return-from-trap** instruction returns to program and reduces privilege level to **user mode**

# System Calls

- Classically, special **trap** instruction (Some modern processors offer less heavy methods)
  - Raises privilege level to **kernel mode**
  - Jumps into the kernel code
  - The **return-from-trap** instruction returns to program and reduces privilege level to **user mode**
- The same dispatch mechanism is used on:
  - System call (process needs OS service)
  - Program fault (illegal memory access, division by zero, etc.)
  - Interrupt (event from an external device)

# System Calls

- CPU checks **mode bit** before executing protected instructions
- OS sets up a **trap table** on initialization
  - List of **trap handlers**
  - Tells the hardware what code should run for various events (e.g., **system call**)

# System Calls

- The Decode stage (Fetch-Decode-Execute) changes:
  - Protected instruction and not kernel mode? raise a **trap**
  - Usually the trap handler will terminate the process
- There are multiple modes: **protection rings**
  - Kernel, executive, superviso, user, etc...
  - Linux uses kernel and user modes only.

# System Calls

- But system calls look like function calls!
  - It is a function call! into the C library
  - But inside is the **trap** instruction
  - Put arguments and system-call number in well-known locations, then execute the trap
  - Unpack return values, return control to the program

# System Calls

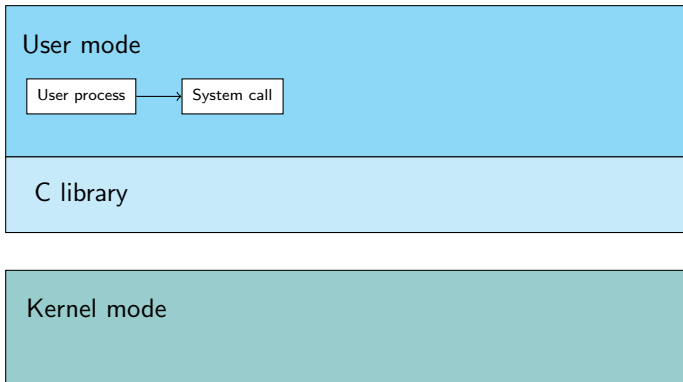




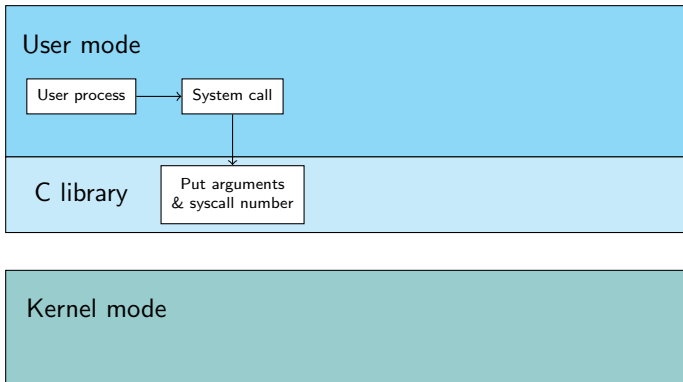
# System Calls



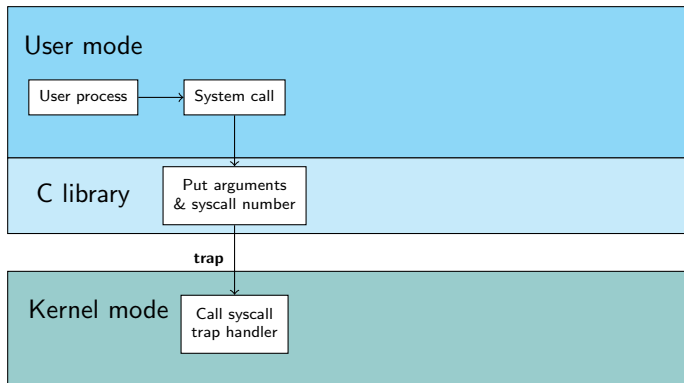
# System Calls



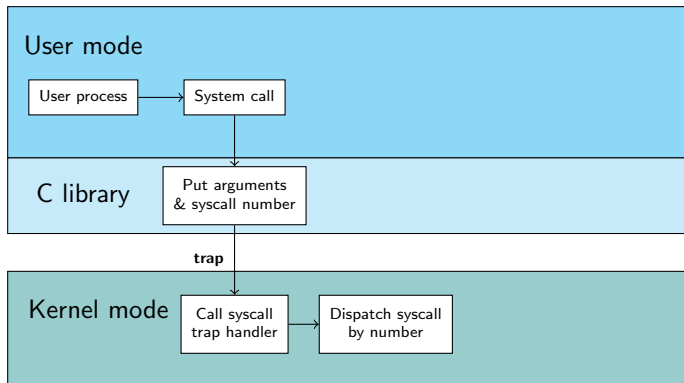
# System Calls



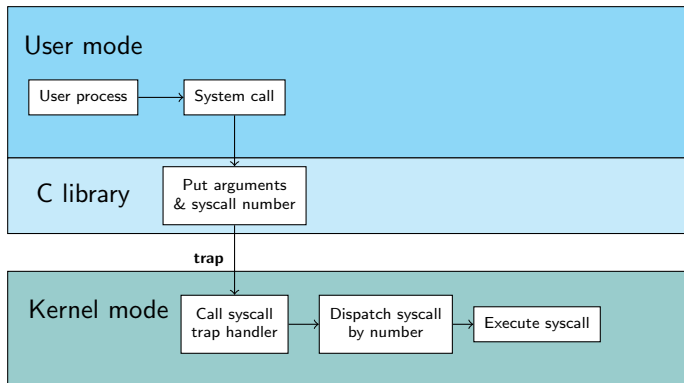
# System Calls



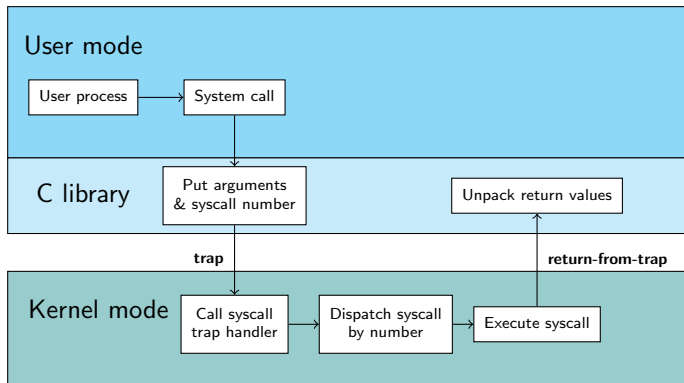
# System Calls



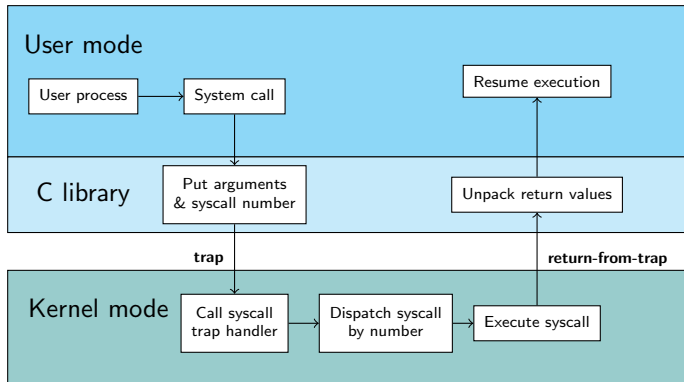
# System Calls



# System Calls

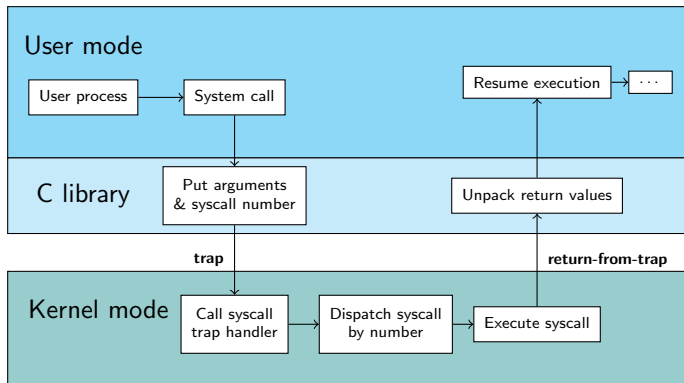


# System Calls





# System Calls



# Limited Execution

Two problems:

- Full access to hardware
  - A bug could corrupt the hardware, freeze the machine, etc.
- How does the OS stop a running process?
  - Implementing **time sharing**
  - What if a process goes into an infinite loop?

Dealt with problem #1

- User/kernel mode, system calls

# Switching Between Processes

- A process is running on the CPU
- The OS (by definition) is not running
- How can the OS regain control of the CPU (to switch between processes)?

# Switching Between Processes

- The **cooperative** approach
  - Most processes frequently make **system calls**
  - Trust the processes to behave reasonably
  - Include an explicit **yield** system call
- But what if a process goes into an infinite loop?

# Switching Between Processes

- The **cooperative** approach
  - Most processes frequently make **system calls**
  - Trust the processes to behave reasonably
  - Include an explicit **yield** system call
- But what if a process goes into an infinite loop?
  - OS can't do much
  - Only recourse: reboot the machine

# Timer Interrupt

- The **preemptive** approach
  - On initialization, the OS starts the timer
  - Timer device programmed to raise an interrupt regularly
  - The timer **interrupt** routine might decide to:
    - Halts current process
    - Saves state of running process
    - (Easily restored later)

# Context Switch

- The **scheduler** makes a decision:
  - Continue running the currently-running process
  - Switch to a different process

# Context Switch

- The **scheduler** makes a decision:
  - Continue running the currently-running process
  - Switch to a different process
- A **context switch**:
  - Save a few register values for the currently-running process
  - Restore values for the upcoming process
  - A **return-from-trap** instruction:
    - Upcoming process becomes currently-running process



# Context Switch

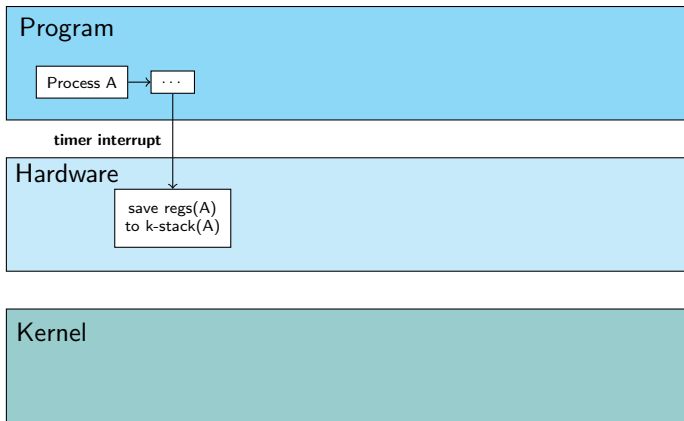
Program

Process A

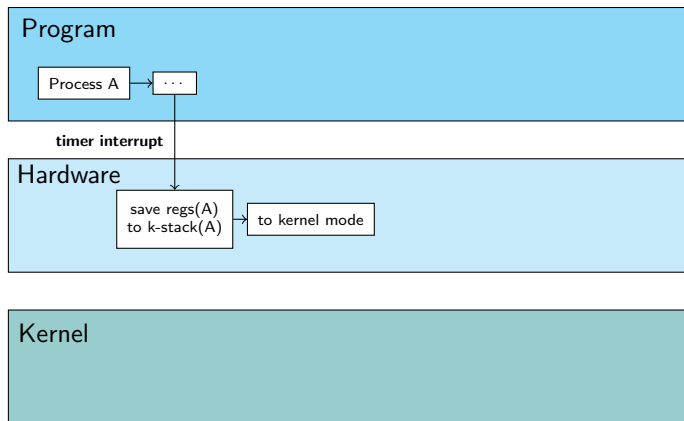
Hardware

Kernel

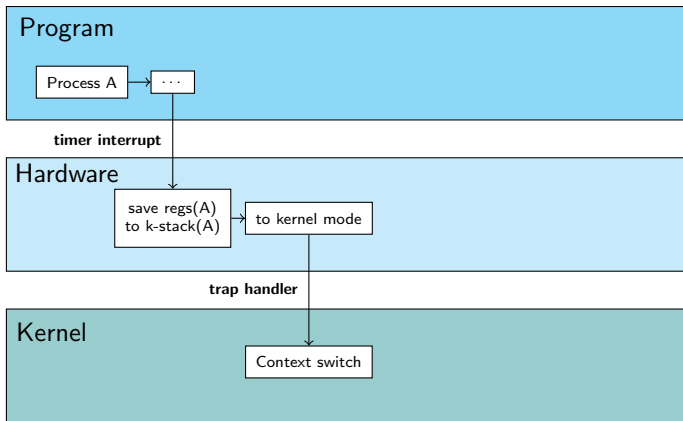
# Context Switch



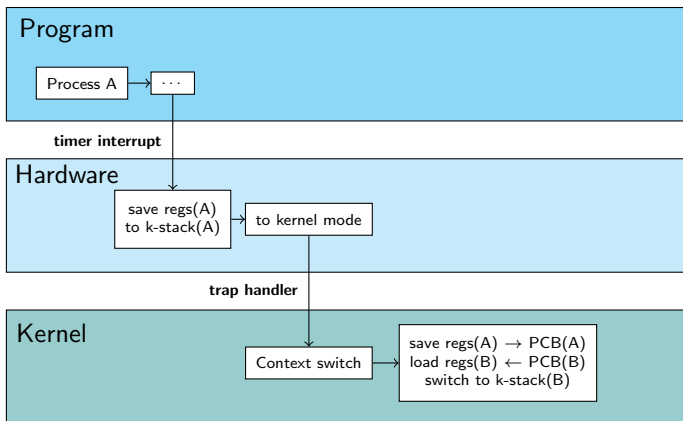
# Context Switch



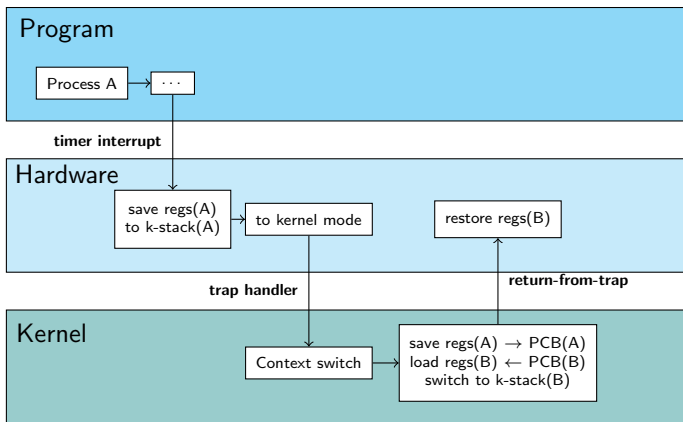
# Context Switch



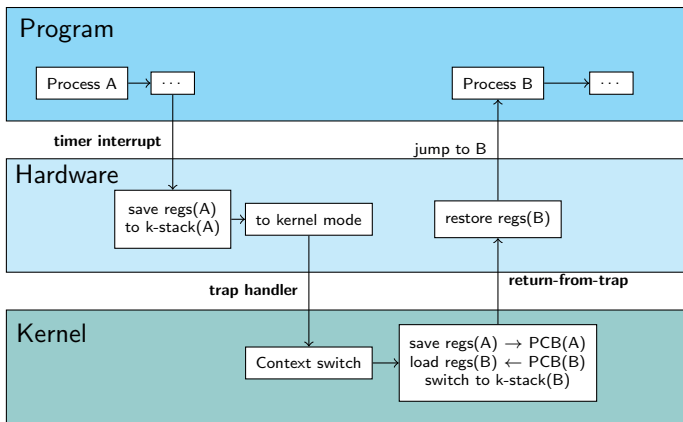
# Context Switch



# Context Switch



# Context Switch



# Context Switch

- A **context switch** should be as fast as possible
  - It is pure overhead
- Context switch is a **mechanism**
- Deciding when to switch (**scheduling**) is a **policy**



# Concurrency

- You should have some concerns...

# Concurrency

- You should have some concerns...
  - Timer interrupt during a system call?
  - Handling one interrupt and another one happens?

# Concurrency

- You should have some concerns...
  - Timer interrupt during a system call?
  - Handling one interrupt and another one happens?
- OS handles these situations:
  - **Disable interrupts** during interrupt processing
    - Too long could lead to lost interrupts
  - make interrupt disabled part **SHORT**
  - Use sophisticated **locking** schemes
    - Protect concurrent access to internal data structures
- More when we get to **concurrency**

# Summary

- **Limited direct execution**

- Run program on the CPU
- Set up hardware to limit what the process can do
- Separate **user mode** and **kernel mode**

- **System calls**

- **Trap** into OS code
- Allow sensitive operations

- **Timer interrupt**

- OS regularly regains control
- **Context switch** when policy dictates