

Locks (ch 28) ptr. 1

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

- **Concurrency** issues
 - Execute a series of instructions **atomically**
 - With interrupts and concurrent processors
- Introducing: a **lock**
 - Critical section seemingly executes atomically

- **Lock variable**

- Holds lock state
- **Available** (or unlocked or **free**)
 - No thread holds the lock
- **Acquired** (or **locked** or **held**)
 - Exactly one thread (**owner**) holds the lock
 - In a critical section

Basic Idea

- `lock()`
 - Try to acquire the lock
 - Will not return (or fail) if held by another thread
- `unlock()`
 - Lock is available again

Basic Idea

- Critical section:

```
balance = balance + 1;
```

- To use lock:

```
1 lock_t mutex; // lock variable
2 ...
3 lock(&mutex);
4 balance = balance + 1;
5 unlock(&mutex);
```

Pthread Locks

- POSIX library: **mutex** (**mutual exclusion**)
- Equivalent code:

```
1 pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2 ...  
3 pthread_mutex_lock(&lock); // may fail!  
4 balance = balance + 1;  
5 pthread_mutex_unlock(&lock);
```

Pthread Locks

- POSIX library: **mutex** (**mutual exclusion**)
- Equivalent code:

```
1 pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
2 ...  
3 pthread_mutex_lock(&lock); // may fail!  
4 balance = balance + 1;  
5 pthread_mutex_unlock(&lock);
```

- Variable passed to lock and unlock
 - May use different locks for different sections
 - **Coarse-grained** locking: one big lock
 - **Fine-grained**: use various locks for different sections

Building A Lock

- Efficient locks provide mutual exclusion at low cost (overhead)
 - Support from hardware and the OS

How can we build an efficient lock?

Evaluating Locks

- **Mutual exclusion**

- At most one thread in the CS

- **Deadlock-freedom**

- Some thread eventually enters CS

- **Fairness (starvation-freedom)**

- Each thread eventually enters CS

- **Performance**

- Time overhead for using the lock
 - Single thread: overhead for grab & release
 - Multiple threads and CPUs

Controlling Interrupts

- Early solution: disable interrupts
 - For single-processor systems
 - No clock interrupt / context switch in critical section
- The negatives:

Controlling Interrupts

- Early solution: disable interrupts
 - For single-processor systems
 - No clock interrupt / context switch in critical section
- The negatives:
 - Trust arbitrary (greedy, malicious, or faulty) programs
 - Does not work on multiprocessors
 - Lost interrupts
- Used by OS

Just Using Loads/Stores

```
1 typedef struct __lock_t { int flag; } lock_t;
2
3 void init(lock_t* mutex) {
4     // 0: available, 1: locked
5     mutex->flag = 0;
6 }
7 void lock(lock_t* mutex) {
8     while (mutex->flag == 1)
9         ; // spin-wait
10    mutex->flag = 1;
11 }
12 void unlock(lock_t* mutex) {
13     mutex->flag = 0;
14 }
```

Just Using Loads/Stores

- No mutual exclusion



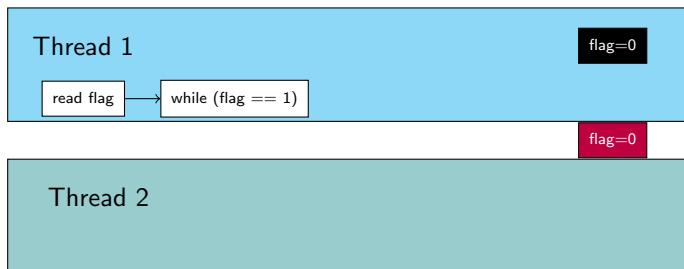
Just Using Loads/Stores

- No mutual exclusion



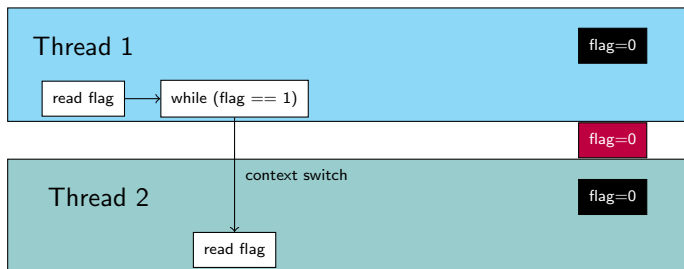
Just Using Loads/Stores

- No mutual exclusion



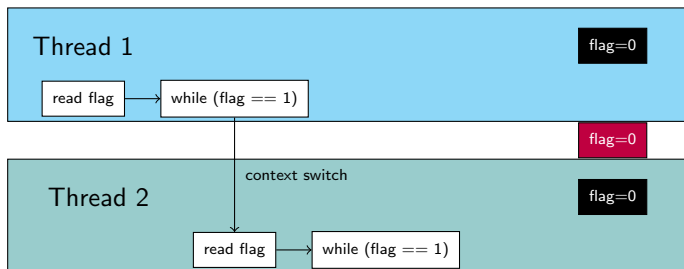
Just Using Loads/Stores

- No mutual exclusion



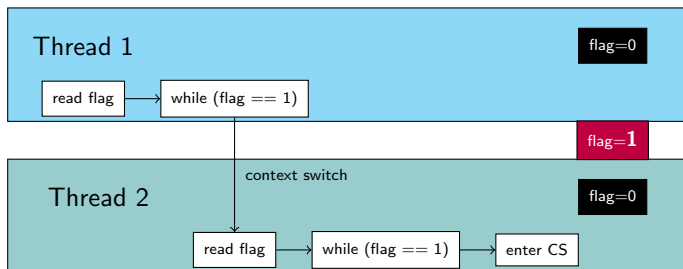
Just Using Loads/Stores

- No mutual exclusion



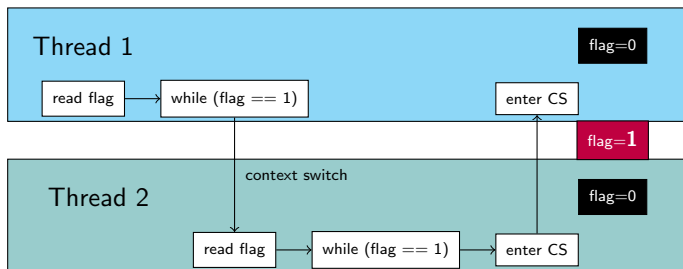
Just Using Loads/Stores

- No mutual exclusion



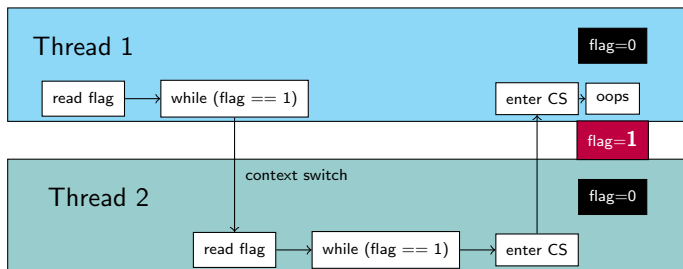
Just Using Loads/Stores

- No mutual exclusion



Just Using Loads/Stores

- No mutual exclusion



Test-And-Set

- Hardware support: a new instruction **test-and-set**
 - Update value and return previous, **atomically**
- Defined as:

```
1 int TestAndSet(int* old_ptr, int new) {  
2     int old = *old_ptr;  
3     *old_ptr = new;  
4     return old;  
5 }
```

New Spin Lock

```
1 typedef struct __lock_t { int flag; } lock_t;
2
3 void init(lock_t* mutex) {
4     // 0: available, 1: locked
5     mutex->flag = 0;
6 }
7 void lock(lock_t* mutex) {
8     while (TestAndSet(&mutex->flag, 1))
9         ; // spin-wait
10    mutex->flag = 1;
11 }
12 void unlock(lock_t* mutex) {
13     mutex->flag = 0;
14 }
```

Evaluating Spin Locks

- Correctness (**mutual exclusion**)?

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom?**

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**?

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**? no
- Performance?
 - Single CPU:

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**? no
- Performance?
 - Single CPU:painful
 - Owner thread is preempted, all $N - 1$ others spin-wait needlessly
 - Multiple CPUs:

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**? no
- Performance?
 - Single CPU: painful
 - Owner thread is preempted, all $N - 1$ others spin-wait needlessly
 - Multiple CPUs: reasonably well

Compare-And-Swap

- Another hardware primitive: **compare-and-swap**
- Compare to `expected`, update only if equal, return previous
- Defined as:

```
1 int CompareAndSwap(int* ptr, int expected, int new) {  
2     int original = *ptr;  
3     if (original == expected)  
4         *ptr = new;  
5     return original;  
6 }
```

Compare-And-Swap

- Spin-lock with CAS:

```
1 void lock(lock_t* lock) {  
2     while (CompareAndSwap(&mutex->flag, 0, 1) == 1)  
3         ; // spin  
4 }
```

- Fairness? performance?

Compare-And-Swap

- Spin-lock with CAS:

```
1 void lock(lock_t* lock) {  
2     while (CompareAndSwap(&mutex->flag, 0, 1) == 1)  
3         ; // spin  
4 }
```

- Fairness? performance?
 - Pretty much the same