

# Synchronization Primitives pt. 2(ch. 30+31+32)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

# Semaphores

- What value should a semaphore be initialized to?
  - General rules: number of resources
  - Lock: 1  $\rightarrow$  can be locked after initialization
  - Ordering: 0  $\rightarrow$  nothing to give away at the start

# Producer / Consumer

- How can we implement a **bounded buffer** with semaphores?

```
1  int buffer[MAX];
2  int fill = 0;
3  int use  = 0;
4
5  void put(int value) {
6      buffer[fill] = value;
7      fill = (fill + 1) % MAX;
8  }
9  int get() {
10     int tmp = buffer[use];
11     use = (use + 1) % MAX;
12     return tmp;
13 }
```

# First Attempt

- What is the problem?

```
1 sem_t empty; // initialized to MAX
2 sem_t full;  // initialized to 0
3
4 void produce(int value) {
5     sem_wait(&empty);
6     put(value);
7     sem_post(&full);
8 }
9 int consume() {
10    sem_wait(&full);
11    int tmp = get();
12    sem_post(&empty);
13    return tmp;
14 }
```

# First Attempt

Producer 1

Producer 2

# First Attempt

Producer 1

```
put(5)
```

Producer 2

# First Attempt

Producer 1

put(5)




buffer[0]=5

Producer 2

# First Attempt

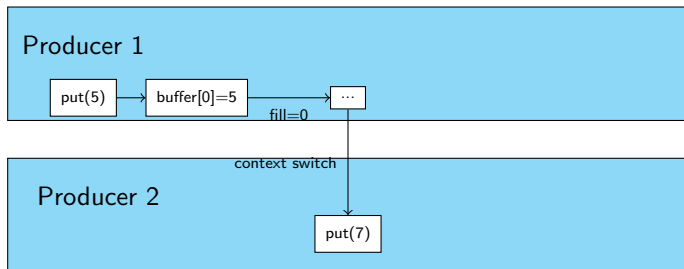
Producer 1




Producer 2




# First Attempt



# First Attempt



# First Attempt



# With Mutual Exclusion

- What is the problem?

```
1  sem_t mutex; // binary semaphore
2  sem_t empty; // initialized to MAX
3  sem_t full;  // initialized to 0
4
5  void produce(int value) {
6      sem_wait(&mutex);
7      sem_wait(&empty);
8      put(value);
9      sem_post(&full);
10     sem_post(&mutex);
11 }
12 int consume() {
13     sem_wait(&mutex);
14     sem_wait(&full);
15     int tmp = get();
16     sem_post(&empty);
17     sem_post(&mutex);
18     return tmp;
19 }
```

# With Mutual Exclusion

- What is the problem? **deadlock!**
  - Producer waits on empty, holds mutex, consumer can't consume

```
1 sem_t mutex; // binary semaphore
2 sem_t empty; // initialized to MAX
3 sem_t full; // initialized to 0
4
5 void produce(int value) {
6     sem_wait(&mutex);
7     sem_wait(&empty);
8     put(value);
9     sem_post(&full);
10    sem_post(&mutex);
11 }
12 int consume() {
13     sem_wait(&mutex);
14     sem_wait(&full);
15     int tmp = get();
16     sem_post(&empty);
17     sem_post(&mutex);
18     return tmp;
19 }
```

# With Mutual Exclusion

- Solution: use mutex around the critical section

```
1 sem_t mutex; // binary semaphore
2 sem_t empty; // initialized to MAX
3 sem_t full; // initialized to 0
4
5 void produce(int value) {
6     sem_wait(&empty);
7     sem_wait(&mutex);
8     put(value);
9     sem_post(&mutex);
10    sem_post(&full);
11 }
12 int consume() {
13     sem_wait(&full);
14     sem_wait(&mutex);
15     int tmp = get();
16     sem_post(&mutex);
17     sem_post(&empty);
18     return tmp;
19 }
```

# Reader-Writer Locks

- More flexible locking primitive
  - e.g., concurrent operations: inserts and lookups
  - Insert changes state → traditional critical section
  - Lookup reads data structure → many at once (if no insert)
- **Reader-writer lock**
  - Four operations: acquire/release read/write lock

# Reader-Writer Locks

- A **single writer** can acquire the lock
- Once a reader acquires a **read lock**:
  - **More readers** are allowed to acquire the read lock
  - A writer waits until all readers are finished



# Reader-Writer Locks

- A **single writer** can acquire the lock
- Once a reader acquires a **read lock**:
  - **More readers** are allowed to acquire the read lock
  - A writer waits until all readers are finished

```
1 typedef struct _rwlock_t {
2     sem_t lock;           // binary semaphore
3     sem_t writelock;      // used to allow ONE writer
4     int readers;          // count of readers in CS
5 } rwlock_t;
6
7 void rwlock_init(rwlock_t* rw) {
8     rw->readers = 0;
9     sem_init(&rw->lock, 0, 1);
10    sem_init(&rw->writelock, 0, 1);
11 }
```

# Reader-Writer Locks

```
1 void rwlock_acquire_writelock(rwlock_t* rw) {
2     sem_wait(&rw->writelock);
3 }
4 void rwlock_release_writelock(rwlock_t* rw) {
5     sem_post(&rw->writelock);
6 }
7
8 void rwlock_acquire_readlock(rwlock_t* rw) {
9     sem_wait(&rw->lock);      // CS for readers
10    rw->readers++;
11    if (rw->readers == 1)
12        sem_wait(&rw->writelock); // first reader grabs writelock
13    sem_post(&rw->lock);
14 }
15 void rwlock_release_readlock(rwlock_t* rw) {
16    sem_wait(&rw->lock);      // CS for readers
17    rw->readers--;
18    if (rw->readers == 0)
19        sem_post(&rw->writelock); // last reader releases writelock
20    sem_post(&rw->lock);
21 }
```

# Reader-Writer Locks


- What is the problem?

# Reader-Writer Locks

- What is the problem? **fairness**
  - Easy to **starve** writer
  - How to prevent readers from starving writers?

# Dining Philosophers

- Five philosophers around a table
  - Single fork between each pair
  - Philosophers **think** and **eat**
  - Two forks to eat (left and right)



# Dining Philosophers

- As code:

```
1 while (1) {  
2     think();  
3     getforks();  
4     eat();  
5     putforks();  
6 }
```

```
1 int left(int p) {  
2     return p;  
3 }  
4 int right(int p) {  
5     return (p - 1) % 5;  
6 }
```

# Dining Philosophers

- Use a semaphore for each fork:

```
1 void get_forks(int p) {  
2     sem_wait(&forks[left(p)]);  
3     sem_wait(&forks[right(p)]);  
4 }  
5 void put_forks(int p) {  
6     sem_post(&forks[left(p)]);  
7     sem_post(&forks[right(p)]);  
8 }
```

- The problem?

# Dining Philosophers

- Use a semaphore for each fork:

```
1 void get_forks(int p) {  
2     sem_wait(&forks[left(p)]);  
3     sem_wait(&forks[right(p)]);  
4 }  
5 void put_forks(int p) {  
6     sem_post(&forks[left(p)]);  
7     sem_post(&forks[right(p)]);  
8 }
```

- The problem? **deadlock!**
  - Each philosopher grabs fork on their left
  - All waiting for their right



# Dining Philosophers

- Solution: break the dependency

```
1 void get_forks(int p) {  
2     if (p == 4) {  
3         sem_wait(&forks[right(p)]);  
4         sem_wait(&forks[left(p)]);  
5     }  
6     else {  
7         sem_wait(&forks[left(p)]);  
8         sem_wait(&forks[right(p)]);  
9     }  
10 }
```

# Thread Throttling

- For example: hundreds of threads work in parallel
- Section of code allocates a lot of memory
  - All threads at the same time → exceeds physical memory
  - Machine will start thrashing (swapping to and from the disk)
- Solution?

# Thread Throttling

- For example: hundreds of threads work in parallel
- Section of code allocates a lot of memory
  - All threads at the same time → exceeds physical memory
  - Machine will start thrashing (swapping to and from the disk)
- Solution? **semaphore!**
  - Initialized to max threads we wish to enter code section
  - Surrounds code section, limits concurrent threads in it

# Implementing Semaphores

- Doesn't maintain invariant: negative value  $\rightarrow$  # waiting threads
  - Easier, matches the Linux implementation

```
1  typedef struct __zem_t {
2      int value;
3      pthread_cond_t cond;
4      pthread_mutex_t lock;
5  };
6
7  void zem_init(zem_t* s, int value) {
8      s->value = value;
9      pthread_cond_init(&s->cond);
10     pthread_mutex_init(&s->lock);
11 }
12 void zem_wait(zem_t* s) {
13     pthread_mutex_lock(&s->lock);
14     while (s->value <= 0)
15         pthread_cond_wait(&s->cond, &s->lock);
16     s->value--;
17     pthread_mutex_unlock(&s->lock);
18 }
19 void zem_post(zem_t* s) {
20     pthread_mutex_lock(&s->lock);
21     s->value++;
22     pthread_cond_signal(&s->cond);
23     pthread_mutex_unlock(&s->lock);
24 }
```

# Summary

- **Condition variables**

- Thread waits until a certain condition
- `wait()`, `signal()`
- **Hold lock** while signaling
- Check value **in a loop**

- **Semaphore**

- Integer value
- Decrement on acquire, wait if negative, increment on release

- **Read-write lock**

- **Single writer** or **multiple readers**

- Dining philosophers

- Think and eat

- Producer / consumer (bounded buffer)