

Introduction (ch. 2) part 2

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

Concurrency

- OS is working on many things at once:
 - Each program has “its own” CPU and RAM
 - Many programs run at the same time
- Modern **multi-threaded** programs exhibit the same problems
- Concurrency is everywhere

Concurrency

threads.c:

```
1  volatile int counter = 0;
2  int loops;
3
4  void* worker(void *arg) {
5      for (int i = 0; i < loops; ++i)
6          ++counter;
7      return NULL;
8  }
9
10 int main(int argc, char *argv[]) {
11     loops = atoi(argv[1]);
12     pthread_t p1, p2;
13     printf("Initial value : %d\n", counter);
14     pthread_create(&p1, NULL, worker, NULL);
15     pthread_create(&p2, NULL, worker, NULL);
16     pthread_join(p1, NULL);
17     pthread_join(p2, NULL);
18     printf("Final value   : %d\n", counter);
19 }
```

Concurrency

- The program creates two **threads**
 - Thread: a function running concurrently within the same address space
- Each thread executes `worker()`
 - Increments a global (shared) counter
- `loops`: how many times to increment the counter

Concurrency

loops = 1000:

```
prompt> gcc -o thread thread.c -Wall -pthread
prompt> ./thread 1000
Initial value : 0
Final value   : 2000
```

loops = 1000000000:

```
prompt> ./thread 1000000000
Initial value : 0
Final value   : 196445738 // huh??
prompt> ./thread 1000000000
Initial value : 0
Final value   : 197944967 // what the??
```

Concurrency

- Three instructions to increment a counter:
 - Load the value into a register
 - Increment the register
 - Store the register back into memory
- Does not execute **atomically** (without interference)
- A problem of **concurrency**

Persistence

- System memory (such as DRAM) is **volatile**
 - Data is lost when power goes away or the system crashes
- We need hardware and software to store data **persistently**

Persistence

- Hardware: **I/O device** such as a **hard drive** or **SSD**
- Software: **file system**
 - Manages the disk
 - Responsible for storing user files
- No private, virtualized disk
- Files can be viewed as virtualized disks.
 - Users want to **share** information that is in files

Persistence

```
1  int main(int argc, char *argv[]) {  
2      int fd = open("/tmp/file", O_WRONLY|O_CREAT,  
3          S_IRWXU);  
4      assert(fd > -1);  
5      int rc = write(fd, "hello world\n", 12);  
6      assert(rc == 12);  
7      close(fd);  
8  }
```

- The program makes three calls:
 - `open()`: opens (and creates) the file
 - `write()`: writes data to the file
 - `close()`: closes the file

Persistence

- These are **System calls**
 - Routed to part of the OS called the **file system**
 - Handles requests and returns error code
 - Like a **standard library** for OS operations
- The **file system**:
 - Figures out where on disk the new data will reside
 - Updates various structures
 - Issues I/O requests to the underlying storage device

- The **kernel** is a core part of the OS:
 - Always in memory
 - Executed in response to events
 - External events (**interrupts**), e.g., clock
 - Requests from running programs (**system calls**)
- What we think of as “OS” is not always part of the kernel
 - e.g., the Unix Shell is an application
- The kernel is an **event driven** program

Kernel - Event Handler


- An **event**: mouse is moved, key is pressed, network communication, division by zero, system call, etc.
- **Interrupts** the current program, executes kernel code
- The kernel defines a **handler** for each event type

- **Interrupt** - asynchronous (external) event
 - For example: key pressed
 - Kernel stores it, can be checked later
- **Trap** - synchronous (internal) event
 - Also **exception** or **fault**
 - For example: division by zero, kernel terminates program
 - Not necessarily for errors!

System Calls


- User program wants to invoke OS - places a **system call**
 - For example: open a file, allocate memory, get keyboard input
- Special instruction that causes a **trap**
- Calls a procedure in the **kernel**
 - The specific **event handler**

Physical Memory




- The kernel resides in memory

Physical Memory




- The code & data of each running program (a **process**) is loaded into memory
 - RAM is divided: user, kernel
 - **System call**: invoke kernel code, then return to user code

Physical Memory



- Several processes can exist in parallel
 - **Memory protection:** each process is seemingly alone

Physical Memory



- Several processes can exist in parallel
 - **Memory protection:** each process is seemingly alone
- This is a major simplification
 - But it suffices for now

Design Goals

- Abstraction
 - Dividing into small, understandable pieces
 - Make the system convenient and easy to use
- Performance
 - Minimize the overhead of the OS
 - Provide virtualization without excessive overheads
- Protection
 - Malicious or bad behavior of one application does not harm others or the OS
 - **Isolation** of processes
- Reliability
 - The OS must run non-stop

Design Goals

- Other goals:
 - Energy efficiency
 - Security
 - Mobility

Some History

- Early OS: Just libraries
 - Commonly-used functions, e.g., low-level I/O
 - No abstraction, no virtualization
 - One program at a time
 - **Batch mode**

Some History

- Early OS: Just libraries
 - Commonly-used functions, e.g., low-level I/O
 - No abstraction, no virtualization
 - One program at a time
 - **Batch mode**
- Beyond Libraries: Protection
 - **User mode** with hardware restrictions
 - **System call**: instead of a library procedure
 - Raises privilege to **kernel mode**
 - OS has full access to hardware

Some History

- Era of Multiprogramming
 - Make better use of machine resources
 - Load a number of jobs, switch rapidly between them
 - **Context switch, concurrency**
 - **Memory protection** became important

Some History

- Era of Multiprogramming
 - Make better use of machine resources
 - Load a number of jobs, switch rapidly between them
 - **Context switch, concurrency**
 - **Memory protection** became important
- The Modern Era
 - The **PC**: the dominant force in computing

Summary

- The OS: abstraction & resource management
- Multiprogramming & concurrency via context switching and memory protection
- Kernel: OS code & data that is always in memory
 - Not a running program, but pieces of code executed in response to events
- System calls: user events to trigger kernel code to act on their behalf