

Locks (ch. 28) pt. 2

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

User mode

Pthread Locks

- POSIX library: **mutex** (mutual exclusion)

```
1 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;  
2 ...  
3 pthread_mutex_lock(&mutex); // may fail!  
4 balance = balance + 1;  
5 pthread_mutex_unlock(&mutex);
```

Pthread Locks

- POSIX library: **mutex** (mutual exclusion)

```
1 pthread_mutex_t mutex = PTHREAD_MUTEX_INITIALIZER;  
2 ...  
3 pthread_mutex_lock(&mutex); // may fail!  
4 balance = balance + 1;  
5 pthread_mutex_unlock(&mutex);
```

- Variable passed to lock and unlock
 - May use different locks for different sections
 - **Coarse-grained** locking: one big lock
 - **Fine-grained**: use various locks for different sections

Building A Lock

- Efficient locks provide mutual exclusion at low cost (overhead)
 - Support from hardware and the OS

How can we build an efficient lock?

Issues compared to kernel

- Can not disable interrupts
- Can not use busy-waits
 - Scheduling out while in busy-wait is catastrophic.
- So, what do we do? We ask for the kernel for help

yield

- Assume there is a yield system call.

```
1 lock(int *mutex) {  
2     while (tas(mutex,1)) yield();  
3 }  
4  
5 unlock(int *mutex) {  
6     *mutex = 0;  
7 }
```

Good?

- It works.
- We depend on the scheduler.
- Quite a lot of system calls (yield) might occur.
- So? If we call the kernel, lets call it for actual work.

Kernel support for mutexes

Three system calls:

```
1 int handle = createMutex();  
2  
3 lockMutex(handle);  
4  
5 unlockMutex(handle);
```

Kernel Implementation: Data structure

```
1 struct userMutex {  
2     int lock;  
3     int hardMutex;  
4     struct queue *queue;  
5 }
```

Kernel Implementation: Lock syscall

```
1 sysLock(struct userMutex *userMutex) {  
2     hardLock(&userMutex->hardMutex);  
3     if (userMutex->lock == 1) {  
4         queueAdd(&userMutex->queue, proc);  
5         BLOCK(&userMutex->hardMutex);  
6     }  
7     userMutex->lock = 1;  
8     hardUnlock(&userMutex->hardMutex);  
9 }
```

Kernel Implementation: unlock syscall

```
1 sysUnock(struct userMutex *userMutex) {  
2     hardLock(&userMutex->hardMutex);  
3  
4     if (userMutex->queue != NULL)  
5         int p = queueRemove(&userMutex->  
6             queue;  
7             UNBLOCK(p);  
8     }  
9     userMutex->lock = 0;  
10    hardUnlock(&userMutex->hardMutex);  
11 }
```

Good?

- It works.
- (For many years this was the standard way)
- What is the problem?
 - Many switches to the kernel.
 - This is considerable overhead for nowadays applications.
 - Most of the time there is no lock contention.
 - Can we exploit this phenomenon?

- Linux: **futex**
 - `futex_wait(address, expected)`
 - Puts calling thread to sleep if address is equal to expected
 - `futex_wake(address)`
 - Wakes one thread waiting on address

Using Queues: Different OS

- Snippet from POSIX thread library:

```
1 void mutex_lock(int *mutex) {
2     int v;
3     // Bit 31 was clear, we got the mutex (fastpath)
4     if (atomic_bit_test_set(mutex, 31) == 0)
5         return;
6     atomic_increment(mutex);
7     while (1) {
8         if (atomic_bit_test_set(mutex, 31) == 0) {
9             atomic_decrement(mutex);
10            return;
11        }
12        v = *mutex;
13        if (v >= 0)
14            continue;
15        futex_wait(mutex, v);
16    }
17 }
18 void mutex_unlock(int *mutex) {
19     if (atomic_add_zero(mutex, 0x80000000))
20         return; // zero iff no other interested threads
21
22     // there are other threads waiting
23     futex_wake(mutex);
24 }
```

Good?

- It works.
- Moreover, if there is no contention then we stay in user mode!

Just Using Loads/Stores

- No mutual exclusion




Just Using Loads/Stores

- No mutual exclusion




Just Using Loads/Stores

- No mutual exclusion



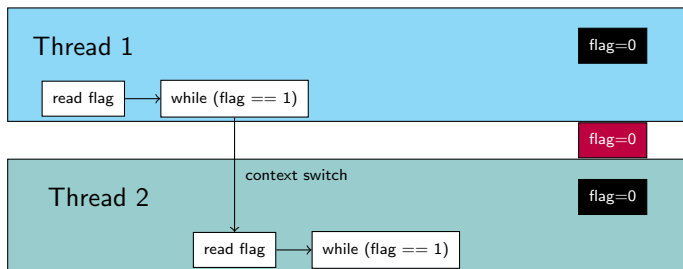
Just Using Loads/Stores

- No mutual exclusion




Just Using Loads/Stores

- No mutual exclusion




Just Using Loads/Stores

- No mutual exclusion



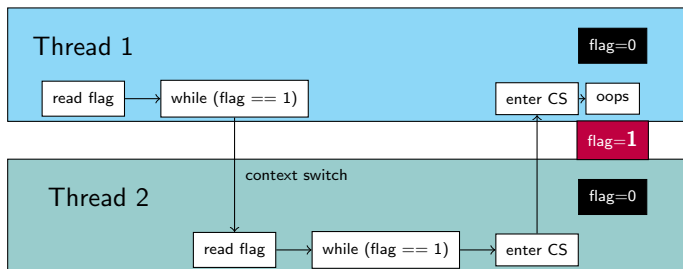
Just Using Loads/Stores

- No mutual exclusion



Just Using Loads/Stores

- No mutual exclusion



Evaluating Spin Locks

- Correctness (**mutual exclusion**)?

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom?**

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**?

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**? no
- Performance?
 - Single CPU:

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**? no
- Performance?
 - Single CPU:painful
 - Owner thread is preempted, all $N - 1$ others spin-wait needlessly
 - Multiple CPUs:

Evaluating Spin Locks

- Correctness (**mutual exclusion**)? yes
- **Deadlock-freedom**? yes
- **Fairness**? no
- Performance?
 - Single CPU: painful
 - Owner thread is preempted, all $N - 1$ others spin-wait needlessly
 - Multiple CPUs: Might be reasonably well

Fetch-And-Add

- Final hardware primitive: **fetch-and-add**
- Atomically increment a value and return old value
- Defined as:


```
1 int FetchAndAdd(int* ptr) {  
2     int old = *ptr;  
3     *ptr = old + 1;  
4     return old;  
5 }
```

Fetch-And-Add

- We can now build a fair ticket lock:

```
1  typedef struct __lock_t {
2      int ticket;
3      int turn;
4  } lock_t;
5
6  void init(lock_t* lock) {
7      lock->ticket = 0;
8      lock->turn = 0;
9  }
10 void lock(lock_t* lock) {
11     int myturn = FetchAndAdd(&lock->ticket);
12     while (lock->turn != myturn)
13         ; // spin
14 }
15 void unlock(lock_t* lock) {
16     lock->turn = lock->turn + 1;
17 }
```



Ticket Lock



Ticket
Turn

Ticket Lock


- A: lock(), gets ticket 0 & runs



Ticket
Turn

Ticket Lock


- A: lock(), gets ticket 0 & runs
- B: lock(), gets ticket 1, spins



Ticket
Turn


Ticket Lock

- A: lock(), gets ticket 0 & runs
- B: lock(), gets ticket 1, spins
- C: lock(), gets ticket 2, spins



Ticket Lock


- A: lock(), gets ticket 0 & runs
- B: lock(), gets ticket 1, spins
- C: lock(), gets ticket 2, spins
- A: unlock(), turn++, B runs



Ticket
Turn


Ticket Lock

- A: lock(), gets ticket 0 & runs
- B: lock(), gets ticket 1, spins
- C: lock(), gets ticket 2, spins
- A: unlock(), turn++, B runs
- A: lock(), gets ticket 3, spins



Ticket Lock

- A: lock(), gets ticket 0 & runs
- B: lock(), gets ticket 1, spins
- C: lock(), gets ticket 2, spins
- A: unlock(), turn++, B runs
- A: lock(), gets ticket 3, spins
- B: unlock(), turn++, C runs
- ...



Two-Phase Locks

- Hybrid approach: **two-phase lock**
 - Spinning can be useful
 - Particularly if lock is about to be released
- **First phase:** lock spins for a while
- **Second phase:** caller put to sleep, wakes up when lock becomes free

Summary

- **Lock**

- Execute a series of actions **atomically**
- Evaluated by: **Mutual exclusion, Deadlock-freedom, fairness, performance**
- POSIX library: **mutex, futex**
- Disabling interrupts: problematic, used by OS
- Hardware support: **test&set, compare&swap, fetch&add**
- Spin-locks: TAS lock & CAS lock
 - Avoid spinning with `yield()`
- Fairness: **ticket lock** or queue lock
- **Condition variables**