

Synchronization Primitives pt 1. (ch. 30+31+32)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

Condition Variables

- Many cases a thread wishes to wait until a certain **condition**
- e.g., waiting for another thread to complete
 - Often called a `join()`
- Shared variable: works, but hugely inefficient

How should a thread wait for a condition?

Condition Variables

- **Waiting** on a condition
 - Thread puts itself in a queue until some state of execution
- **Signaling** on a condition
 - Some other thread can wake waiting thread

Condition Variables

- **Waiting** on a condition
 - Thread puts itself in a queue until some state of execution
- **Signaling** on a condition
 - Some other thread can wake waiting thread
- Name is a bit misleading
 - More of a queue
 - We are responsible for the actual “condition”

Definitions and Routines

- Declare a condition variable:

```
pthread_cond_t cv;
```

- Operations:

```
// wait:  
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);  
// signal:  
pthread_cond_signal(pthread_cond_t *c);
```

- Wait call takes **mutex** as a parameter
 - Caller must be its **owner** (have it locked)
 - Releases the lock, puts caller to sleep
 - On wake up, re-acquires lock and returns

Parent Waiting For Child

- Two threads:
 - **Parent:**
 - Creates child thread
 - Waits on CV until child completes
 - **Child:**
 - Prints a message (“child”)
 - Wakes parent by signaling on CV

Parent Waiting For Child

```
1 pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;
2 pthread_cond_t c = PTHREAD_COND_INITIALIZER;
3
4 void* child(void* arg) {
5     printf("child\n");
6     thr_exit();
7     return NULL;
8 }
9 int main(void) {
10    printf("parent: begin\n");
11    pthread_t p;
12    pthread_create(&p, NULL, child, NULL);
13    thr_join();
14    printf("parent: end\n");
15    return 0;
16 }
```

Parent Waiting For Child

```
1 void thr_exit() {  
2     pthread_cond_signal(&c);  
3 }  
4 void thr_join() {  
5     pthread_mutex_lock(&m);  
6     pthread_cond_wait(&c, &m);  
7     pthread_mutex_unlock(&m);  
8 }
```

- Why might this code fail?

Parent Waiting For Child

```
1 void thr_exit() {  
2     pthread_cond_signal(&c);  
3 }  
4 void thr_join() {  
5     pthread_mutex_lock(&m);  
6     pthread_cond_wait(&c, &m);  
7     pthread_mutex_unlock(&m);  
8 }
```

- Why might this code fail?
 - Child runs immediately
 - Will signal, but no thread asleep on CV
 - Parent runs, calls wait and gets stuck
 - Solution?

Parent Waiting For Child

```
1 void thr_exit() {  
2     pthread_cond_signal(&c);  
3 }  
4 void thr_join() {  
5     pthread_mutex_lock(&m);  
6     pthread_cond_wait(&c, &m);  
7     pthread_mutex_unlock(&m);  
8 }
```

- Why might this code fail?
 - Child runs immediately
 - Will signal, but no thread asleep on CV
 - Parent runs, calls wait and gets stuck
 - Solution? use done variable

Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      done = 1;
4      pthread_cond_signal(&c);
5  }
6  void thr_join() {
7      pthread_mutex_lock(&m);
8      if (done == 0)
9          pthread_cond_wait(&c, &m);
10     pthread_mutex_unlock(&m);
11 }
```

- Why might this code fail?

Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      done = 1;
4      pthread_cond_signal(&c);
5  }
6  void thr_join() {
7      pthread_mutex_lock(&m);
8      if (done == 0)
9          pthread_cond_wait(&c, &m);
10     pthread_mutex_unlock(&m);
11 }
```

- Why might this code fail?
 - Parent calls join, sees done=0
 - Interrupted just before wait, context switch to child
 - Child sets done, signal is lost, parent is stuck again
 - Solution?

Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      done = 1;
4      pthread_cond_signal(&c);
5  }
6  void thr_join() {
7      pthread_mutex_lock(&m);
8      if (done == 0)
9          pthread_cond_wait(&c, &m);
10     pthread_mutex_unlock(&m);
11 }
```

- Why might this code fail?
 - Parent calls join, sees done=0
 - Interrupted just before wait, context switch to child
 - Child sets done, signal is lost, parent is stuck again
 - Solution? hold lock while signaling

Parent Waiting For Child

```
1  int done = 0;
2  void thr_exit() {
3      pthread_mutex_lock(&m);
4      done = 1;
5      pthread_cond_signal(&c);
6      pthread_mutex_unlock(&m);
7  }
8  void thr_join() {
9      pthread_mutex_lock(&m);
10     while (done == 0)
11         pthread_cond_wait(&c, &m);
12     pthread_mutex_unlock(&m);
13 }
```

- Additionally, check variable in a loop
 - Condition variable may signal unexpectedly
 - Also crucial for more than 2 threads

Covering Conditions

- Memory allocator implementation
- Assume zero bytes are free:
 - Thread A calls `allocate(100)`
 - Thread B calls `allocate(10)`
 - Both A and B wait on the condition
 - Thread C calls `free(50)`

Covering Conditions

- Memory allocator implementation
- Assume zero bytes are free:
 - Thread A calls `allocate(100)`
 - Thread B calls `allocate(10)`
 - Both A and B wait on the condition
 - Thread C calls `free(50)`
 - Which waiting thread wakes up?

Covering Conditions

- Memory allocator implementation
- Assume zero bytes are free:
 - Thread A calls `allocate(100)`
 - Thread B calls `allocate(10)`
 - Both A and B wait on the condition
 - Thread C calls `free(50)`
 - Which waiting thread wakes up?
- Solution: wake up all waiting threads
 - Use `pthread_cond_broadcast()`
 - Performance cost: too many threads might be woken

Producer / Consumer

- Also: **bounded buffer** problem
- **Producer**
 - Produces data items
 - Wishes to place items in a buffer
- **Consumer**
 - Grabs items out of the buffer
 - Consumes items in some way
- e.g., web server consumes HTTP requests (work queue)

Producer / Consumer

- Also used in pipes:
 - `grep foo file.txt | wc -l`
 - `grep` - output lines from `file.txt` containing `foo`
 - `wc -l` - output number of lines from input
 - Shell redirects `grep` standard output to a **pipe**
 - Created by the `pipe` system call
 - Other end connected to standard input of `wc`
 - `grep` - producer, `wc` - consumer
- In-kernel bounded buffer

Producer / Consumer

```
1 int buffer;
2 int count = 0; // initially, empty
3
4 void put(int value) {
5     assert(count == 0);
6     count = 1;
7     buffer = value;
8 }
9 int get() {
10    assert(count == 1);
11    count = 0;
12    return buffer;
13 }
```

First Attempt

- What are the two problems?

```
1 pthread_cond_t cond;
2 pthread_mutex_t mutex;
3
4 void produce(int i) {
5     pthread_mutex_lock(&mutex);
6     if (count == 1)
7         pthread_cond_wait(&cond, &mutex);
8     put(i);
9     pthread_cond_signal(&cond);
10    pthread_mutex_unlock(&mutex);
11 }
12 int consume() {
13     pthread_mutex_lock(&mutex);
14     if (count == 0)
15         pthread_cond_wait(&cond, &mutex);
16     int tmp = get();
17     pthread_cond_signal(&cond);
18     pthread_mutex_unlock(&mutex);
19     return tmp;
20 }
```

First Attempt

- More than one consumer / producer

First Attempt

- More than one consumer / producer
- Two consumers:
 - No while on CV in consume()
 - Can consume when empty!
 - (and the same for produce())

```
1 int consume() {  
2     pthread_mutex_lock(&mutex);  
3     while (count == 0)  
4         pthread_cond_wait(&cond, &mutex);  
5     int tmp = get();  
6     pthread_cond_signal(&cond);  
7     pthread_mutex_unlock(&mutex);  
8     return tmp;  
9 }
```

First Attempt

buffer=0

Producer

Consumer 1

Consumer 2

First Attempt

buffer=0

Producer

Consumer 1

consume

Consumer 2

First Attempt

buffer=0

Producer

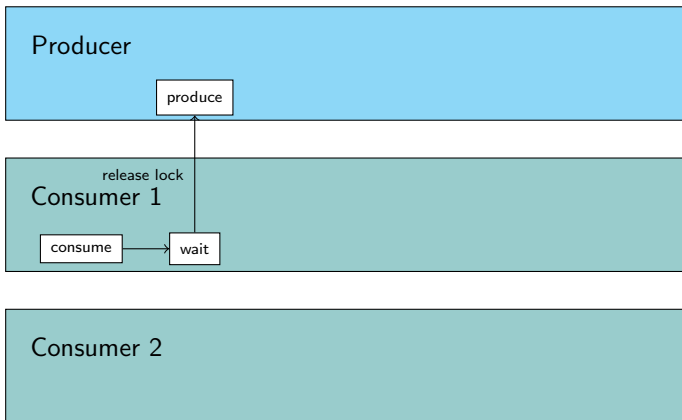
Consumer 1



Consumer 2

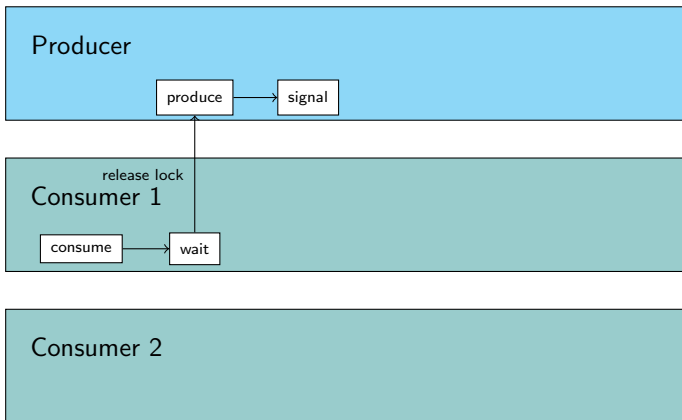
First Attempt

buffer=1



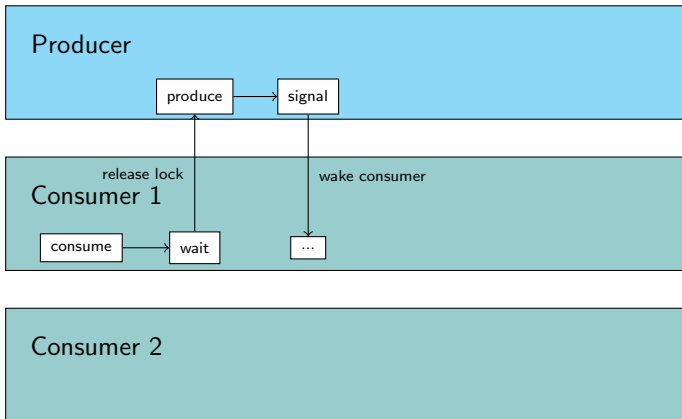
First Attempt

buffer=1



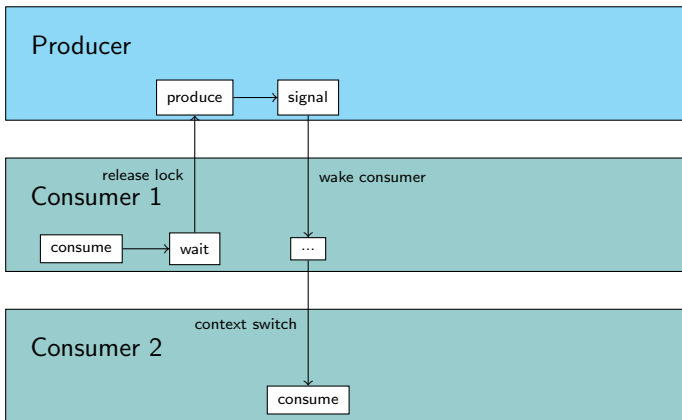
First Attempt

buffer=1



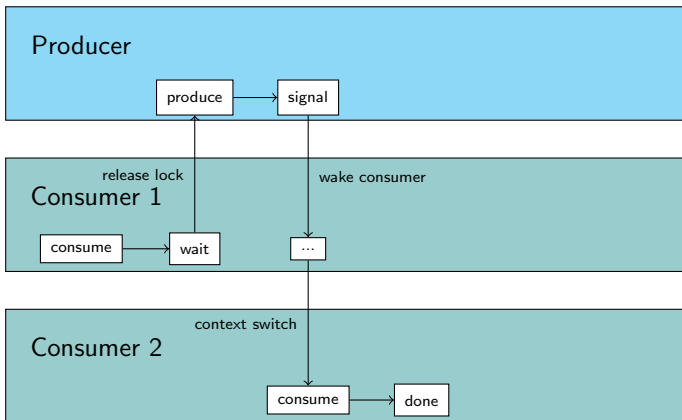
First Attempt

buffer=1



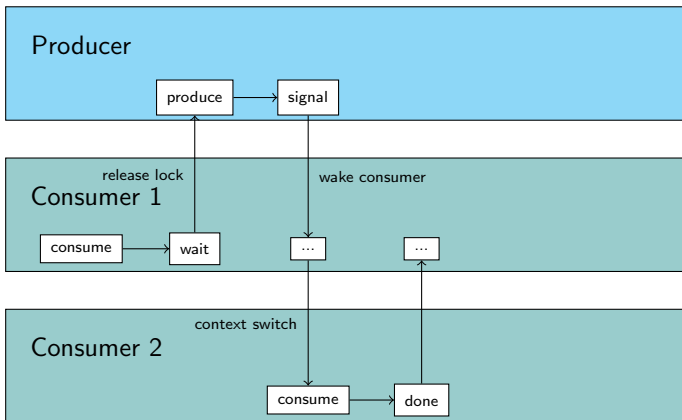
First Attempt

buffer=0



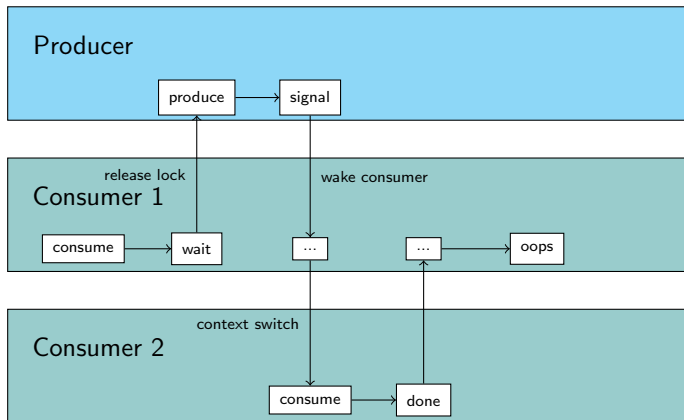
First Attempt

buffer=0



First Attempt

buffer=0



First Attempt

- Only one condition variable!
 - Consumer might wake another consumer
 - Producer might wake another producer
- Solution?

First Attempt

- Only one condition variable!
 - Consumer might wake another consumer
 - Producer might wake another producer
- Solution? use **two** condition variables
 - Producer threads wait on `empty`
 - Consumer threads wait on `full`