

Introduction (ch. 2) part 1

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy


Tel-Aviv Academic College

What is an Operating System?

Not a simple question.

What is an Operating System?

Not a simple question.



- Middleware between user programs and hardware
- Abstracts and manages resources
 - CPU
 - Main memory
 - I/O Devices (disk, network card, mouse, keyboard, monitor, etc.)

Why study Operating Systems?

- We study Computer **Science**
 - Not a programming course...
- You use an operating system
 - The machine is (mostly) useless without an OS
 - Understand what you use
 - Why and how is the OS useful?
- Behavior of OS impacts entire machine
 - Understand system performance
 - Useful to know how computers work

Approach

- The course is about ideas and analysis
- We will not build an OS
 - But there will be coding

CPU wise, What happens when a program runs?

- Executes instructions
- The processor:
 - **Fetches** an instruction from memory
 - **Decodes** the instruction
 - **Executes** it
 - Moves on to the **next instruction**
- **Von Neumann** model of computing

What does the OS do?

- Abstraction
 - Virtual resources that correspond to hardware resources
 - Well-defined operations on these resources
 - CPU → Running program (process / thread)
 - Memory → Address space / virtual memory
 - Storage → Files, file system

What does the OS do?

- Resource Management
 - Share resources among running programs
 - Decide who gets how much and when
 - CPU → Who runs next?
 - Memory → Where is data in RAM and when to access it?
 - Storage → Where and how are files stored on disk?

Three Easy Pieces

- **Virtualization**

- As if each program has resources to itself

- **Concurrency**

- Juggling many things at once

- **Persistence**

- Ability to store data beyond termination / computer shutdown

Virtualizing the CPU

cpu.c:

```
1  int main(int argc, char *argv[])
2  {
3      while (1) {
4          spin(1); // returns after 1 second
5          printf("%s\n", argv[1]);
6      }
7  }
```

- The program loops and prints

Virtualizing the CPU

```
prompt> gcc -o cpu cpu.c -Wall
prompt> ./cpu "A"
A
A
A
A
^C
prompt>
```

- Runs forever
- Halt program by pressing "Control-C"

Virtualizing the CPU

```
prompt> ./cpu A & ./cpu B & ./cpu C & ./cpu D &  
[1] 7353  
[2] 7354  
[3] 7355  
[4] 7356  
A  
B  
D  
C  
A  
B  
D  
C  
A  
C  
B  
D  
...
```

- We have **one** processor
- All four seem to be running **at the same time**

Virtualizing the CPU

- OS illusion
 - Single CPU → infinite number of virtual CPUs
 - Programs seemingly run at once

Virtualizing the CPU

- OS illusion
 - Single CPU → infinite number of virtual CPUs
 - Programs seemingly run at once
- Under the hood:
 - Multiple instances are started
 - OS picks one to run (use CPU)
 - After a while, OS kicks it off the CPU
 - Picks another one to run

Virtualizing the CPU

- **Context Switch**

- Running program pauses, another is brought in
- Program does not know when it is context-switched (in or out)
 - Illusion that it is alone on the CPU
 - Fetch-Decode-Execute cycle continues
- Fast and frequent
 - Appears to be running at the same time

Virtualizing Memory

- Physical memory (**RAM**) - array of bytes
 - **Read** (load) - specify address to access data
 - **Write** (store) - also specify data to write
- Program code (instructions) is also in memory!

Virtualizing Memory

mem.c:

```
1  int main(int argc, char *argv[])
2  {
3      int p;
4      printf("(%d) the address of p: %p\n",
5             getpid(), &p);
6      p = atoi(argv[1]);
7      while (1) {
8          spin(1);
9          p = p + 1;
10         printf("(%d) p: %d\n", getpid(), p);
11     }
12 }
```

Virtualizing Memory

```
prompt> ./mem 0
(2134) the address of p: 0x200000
(2134) p: 1
(2134) p: 2
(2134) p: 3
(2134) p: 4
^C
prompt>
```

- Increments and prints every second

Virtualizing Memory

```
prompt> ./mem 0 & ./mem 100 &  
[1] 13526  
[2] 13527  
(13527) the address of p: 0x200000  
(13526) the address of p: 0x200000  
(13527) p: 101  
(13526) p: 1  
(13527) p: 102  
(13526) p: 2  
(13527) p: 103  
(13526) p: 3  
...
```

- Same address, different value
- As if each instance as its own private memory

Virtualizing Memory

- Program address is not physical address
 - It is a **virtual address**
- Each process accesses its own **virtual address space**
 - The OS (with hardware help) maps it onto the physical memory
 - Reference in one running program does not affect the other

Virtualizing Memory

- Program address is not physical address
 - It is a **virtual address**
- Each process accesses its own **virtual address space**
 - The OS (with hardware help) maps it onto the physical memory
 - Reference in one running program does not affect the other
- Each program seemingly has all physical memory to itself
 - No knowledge (or responsibility) of other programs
 - **Memory protection**