

File System Implementation pt. 2

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseaux

Moshe Sulamy

Tel-Aviv Academic College

Directory Organization

- Directory: list of (entry name, inode number) pairs
- Two extra files: “dot” & “dot-dot” for current and parent dirs
 - e.g., `dir` has three files:

inum	reclen	strlen	name
5	12	2	.
2	12	3	..
12	12	4	foo
13	12	4	bar
24	36	28	foobar_is_a_pretty_longname

Free Space Management

- File system must track free inodes and data blocks
- In vsfs: two bitmaps
 - New file: search through bitmap for free inode, allocate it
 - **Pre-allocation**: commonly used
 - Look for and allocate contiguous blocks for file

Access Paths

- Open a file (`/foo/bar`), read it, close it
 - Issue an `open("/foo/bar", O_RDONLY)`
 - **Traverse** pathname to locate desired inode
 - Begin at root: well-known, usually inode 2
 - Read block that contains inode 2
 - Look inside it - read data block to find inode number of `foo`
 - Read inode and data blocks of `foo` to find `bar`
 - Read from the file, repeat:
 - Read `bar` inode to find data block
 - Read data block
 - Write to inode - update access time
 - Close the file

Access Paths

- Writing to a file:
 - Open file (as before)
 - Each write generates five I/Os:
 - 1 Read data bitmap
 - 2 Write updated data bitmap (newly-allocated block to use)
 - 3 Read the inode
 - 4 Write updated inode with new block location
 - 5 Write the actual block itself

Access Paths

- Creating a file:
 - Read inode bitmap
 - Write updated inode bitmap with allocated inode
 - Write inode itself
 - Write data to directory containing the file
 - Read and write directory inode to update it
 - Directory needs to grow? Additional I/O
 - To data bitmap, new directory block

Caching and Buffering

- Simple operations: huge number of I/Os
 - e.g., long pathname can lead to hundreds of reads
 - Just to open a file!

What can a file system do to reduce the costs of many I/Os?

Caching and Buffering

- Simple operations: huge number of I/Os
 - e.g., long pathname can lead to hundreds of reads
 - Just to open a file!

What can a file system do to reduce the costs of many I/Os?

- Use system memory (DRAM) to cache important blocks
 - Early systems used fixed-size cache
 - Static partitioning of memory: can be wasteful
 - Modern systems use **dynamic partitioning**

Caching and Buffering

- Sufficiently large cache: avoid read I/O altogether
- Write traffic has to go to disk
 - Cache does not reduce write I/O

Caching and Buffering

- Sufficiently large cache: avoid read I/O altogether
- Write traffic has to go to disk
 - Cache does not reduce write I/O
- Use **write buffering**
 - Delay writes: **batch** updates to smaller set of I/Os (several updates to inode bitmap)
 - Buffer writes in memory, **schedule** subsequent I/Os
 - **Avoid** writes, e.g., file created and then deleted
- Use `fsync()` to force writes to disk

UNIX File System

- The old UNIX file system:

UNIX File System

- The old UNIX file system:
 - Superblock: volume size, number of inodes, pointer to head of free list of blocks, etc.
 - The inode region for all inodes
 - Data blocks take up most of the disk



UNIX File System

- The old UNIX file system:
 - Superblock: volume size, number of inodes, pointer to head of free list of blocks, etc.
 - The inode region for all inodes
 - Data blocks take up most of the disk




- Simple, supports basic abstractions, easy to use
- Problem: terrible performance (2% of disk bandwidth)

UNIX File System

- Disk treated as random-access memory
 - Expensive positioning costs
 - e.g., data blocks of file far away from its inode
- File system **fragmented**
 - Logically contiguous file → back and forth across the disk
- Block size too small (512 bytes)
 - Bad for data transfer
 - Positioning overhead for each block


UNIX File System

- For example, data block region with four files:



UNIX File System

- For example, data block region with four files:




- Files B & D are deleted:



UNIX File System

- For example, data block region with four files:



- Files B & D are deleted:



- Allocate file E, of size four blocks:




- **Fast File System**

- Design structures and allocation to be “disk aware”
- Keep same API
 - `open()`, `read()`, `write()`, `close()`, etc.
 - Change internal implementation
 - Paved the path for new file system construction

Cylinder Group

- FFS divides disk into **cylinders** and **cylinder groups**
 - In modern file systems: **block groups**
 - e.g., Linux ext2, ext3, and ext4



Cylinder Group

- Use groups to improve seek performance
- e.g., place two files within the same group
- Allocate files and directories within each group

G0	G1	G2	G3	G4	G5	G6	G7	G8	G9
----	----	----	----	----	----	----	----	----	----

Cylinder Group



- Within a single group:
 - Copy of the **super block** (S)
 - For reliability reasons
 - Per-group **inode bitmap** (ib) and **data bitmap** (db)
 - The **inode** and **data block** regions
 - Same as vsfs

Policies

- Keep related stuff together
 - Keep unrelated stuff far apart
- Placement of directories:
 - Find group with low number of allocated directories, high number of free inodes
 - Put directory data and inode in that group
- Placement of files:
 - Allocate data blocks in same group as inode
 - Place files in same group as directory

Policies

- Create 3 dirs (/, /a, /b) and four files (/a/c, /a/d, /a/e, /b/f)

group	inodes	data	group	inodes	data
0	/-----	/-----	0	/-----	/-----
1	acde-----	accddee---	1	a-----	a-----
2	bf-----	bff-----	2	b-----	b-----
3	-----	-----	3	c-----	cc-----
4	-----	-----	4	d-----	dd-----
5	-----	-----	5	e-----	ee-----
6	-----	-----	6	f-----	ff-----
7	-----	-----	7	-----	-----
...			...		
With name locality			No name locality		

Large-File Exception

- General policy: exception for large files
 - Entirely fill block group it is placed within
 - Prevents related files from being placed in group

G0	G1	G2	G3	G4	G5	G6	G7	G8	G9
		0 1 2 3 4							
		5 6 7 8 9							

- For large files: spread chunks across disk
 - Hurts performance, can address by choosing chunk size carefully
 - Reduce overhead by doing more work: **amortization**

G0	G1	G2	G3	G4	G5	G6	G7	G8	G9
8 9		0 1		2 3		4 5		6 7	

Amortization


- For example:
 - Average positioning time: 10 ms, transfer rate: 40 MB/s
 - Goal: Achieve 50% of peak disk performance
 - i.e., 10 ms transferring data for every 10 ms positioning
 - How big does a chunk have to be?

Amortization


- For example:
 - Average positioning time: 10 ms, transfer rate: 40 MB/s
 - Goal: Achieve 50% of peak disk performance
 - i.e., 10 ms transferring data for every 10 ms positioning
 - How big does a chunk have to be?
- 40 MB/s \rightarrow 409.6 KB
 - $(40 \cdot 1024/100)$

Amortization

- For example:
 - Average positioning time: 10 ms, transfer rate: 40 MB/s
 - Goal: Achieve 50% of peak disk performance
 - i.e., 10 ms transferring data for every 10 ms positioning
 - How big does a chunk have to be?
- 40 MB/s \rightarrow 409.6 KB
 - $(40 \cdot 1024 / 100)$



- Internal fragmentation
 - Most files were small (at the time)
 - Use sub-blocks of 512 bytes
- **Parameterization**
 - Skip over every other block
 - Enough time to request next block before it went past disk head



- **Track buffer** prevents two spins to read track
- Also introduced: **long file names** and **symbolic links**

Summary

- Divide disk into blocks
 - Commonly-used size (4KB)
 - Data region for user data, metadata region for inodes
 - Allocation structure (data and inode bitmaps)
 - **Superblock**: information about file system
- Data uses **direct** and **indirect pointers**
 - Multi-level approach: pointer to block of indirect pointers
 - **Extents**: disk pointer plus length
- Access paths: huge number of I/Os
 - Cache with **dynamic partitioning**
 - **Write buffering**
- FFS: using **cylinder groups** and **large-file exception**