

Locks (ch. 28) pt. 1

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

Tel-Aviv Academic College

Kernel mode, Single processor

Evaluating Locks

- **Mutual exclusion**

- At most one thread in the CS

- **Deadlock-freedom**

- Some thread eventually enters CS

- **Fairness (starvation-freedom)**

- Each thread eventually enters CS

- **Performance**

- Time overhead for using the lock
 - Single thread: overhead for grab & release
 - Multiple threads and CPUs

Get Free Page

```
struct run {  
    struct run *next;  
};  
struct run *freelist;  
  
void *kalloc(void)  
{  
    struct run *r;  
  
    r = freelist;  
    if (r)  
        freelist = r->next;  
  
    return (void*)r;  
}
```

Single processor

Recall

- The kernel is a multi threaded app.
- At minimum each process is a thread in the kernel.
- If there are user mode threads, there might be more kernel threads.
- Context switch happens in the kernel.

So. What is the problem with the kalloc?

Single processor

Recall

- The kernel is a multi threaded app.
- At minimum each process is a thread in the kernel.
- If there are user mode threads, there might be more kernel threads.
- Context switch happens in the kernel.

So. What is the problem with the kalloc?

Race condition.

Solution?

Single processor

Recall

- The kernel is a multi threaded app.
- At minimum each process is a thread in the kernel.
- If there are user mode threads, there might be more kernel threads.
- Context switch happens in the kernel.

So. What is the problem with the kalloc?

Race condition.

Solution?

Critical section.

Single processor, Critical section

```
void *kalloc(void)
{
    struct run *r;

    intDsbl();

    r = freelist;
    if (r)
        freelist = r->next;

    intEnbl();

    return (void*)r;
}
```


Is this good?

- It works.
- The interrupts are disabled for a SHORT period of time.

So, it is reasonable. No better solution.

However. What happens if the critical section is one of the following:

- long
- requires interrupt machinery (which anyways means long)

softLock

```
int mutex = 0;

void lock() {
    if (mutex == 0) {
        mutex = 1;
        return;
    };
    //
    // How do we wait for mutex == 0?
    //
}
```

softLock

```
int mutex = 0;

void lock() {
    if (mutex == 0) {
        mutex = 1;
        return;
    };
    while (mutex == 0);
    mutex = 1;
}
```

Problems

- Ugly code.
- Race conditions.
- Let us add `intDsbl/intEnbl`

softLock

```
int mutex = 0;

void lock() {
    intDsbl();
    if (mutex == 0) {
        mutex = 1;
        intEnbl();
        return;
    };
    while (mutex == 0);
    mutex = 1;
    intEnbl();
}
```

softLock DEADLOCK

- If we get into the busy wait, we are stuck forever.
- We are in kernel. We can BLOCK.

BLOCK

- BLOCKs (i.e., set state to blocked) the kernel thread.

```
1      proc->state = BLOCKED;  
2  //  
3      sched();  
4  //
```

- Note that `intDsbl` is in on process and `intEnbl` in the following one.

softLock

```
int mutex = 0;

void lock() {
    intDsbl();
    if (mutex == 0) {
        mutex = 1;
        intEnbl();
        return;
    };
    BLOCK();
    mutex=1; //One UNBLOCKed, OK. Else , bad.
    intEnbl();
}
```


softLock, good one

```
int mutex = 0;

void lock() {
    intDsbl();
    while (mutex == 1) BLOCK();
    mutex = 1;
    intEnbl();
    return;
}
```

UNBLOCK

- Coarse: UNBLOCK moves ALL blocked kernel threads (i.e., processes) to ready.

```
1  For each process if state BLOCK then  
    set to READY
```

- Fine: UNBLOCK(pid)

```
int mutex = 0;

void unlock() {
    intDsbl();
    mutex = 0;
    UNBLOCK();
    intEnbl();
}
```

- UNBLOCK might be too long.
- The intDsbl/intEnbl create a critical section in the soft lock implementation.
- This is a very coarse lock.
- We can use many mutexes in order to get fine grain soft locks.
- We can use UNBLOCK which wakes up only one thread waiting for the specific mutex.

Kernel mode, Multiple processors

The race condition is BACK!

```
void *kalloc(void)
{
    struct run *r;

    intDsbl();

    r = freelist;
    if (r)
        freelist = r->next;

    intEnbl();

    return (void*)r;
}
```

hardLock

```
1 void hardLock(int *mutex) {  
2     intDsbl();  
3     while ( *mutex != 0 );  
4     mutex = 1;  
5     intEnbl();  
6 }
```

Still race condition!!

Test-And-Set

Machine instruction:

TAS mem,new,reg

- Hardware support: a new instruction **test-and-set**
 - Update value and return previous, **atomically** across all processors
- Defined as:

```
1 int TestAndSet(int* mem, int new) {  
2     int old = *mem  
3     *mem = new;  
4     return old;  
5 }
```


hardLock

```
void hardLock(int *mutex) {  
    intDsbl();  
    while (tas(mutex, 1));  
    intEnbl();  
}  
  
void hardUnlock(int *mutex) {  
    *mutex = 0;  
}
```

TAS not unique: Compare-And-Swap

Machine instruction:

CAS mem,expected,new,reg

- Another hardware primitive: **compare-and-swap**
- Compare to `expected`, update only if equal, return previous
- Defined as:

```
1 int CompareAndSwap(int* ptr, int expected, int new) {  
2     int original = *ptr;  
3     if (original == expected)  
4         *ptr = new;  
5     return original;  
6 }
```

hardLock (equivalent to the prev one)

```
1 void hardLock(int *mutex) {  
2     intDsbl();  
3     while (cas(mutex, 0, 1));  
4     intEnbl();  
5 }  
6  
7 void hardUnlock(int *mutex) {  
8     *mutex = 0;  
9 }
```

Decker's Algorithm: Without hardware support

- What about a lock without hardware support?

```
1 int flag[2];      // wants to grab lock?
2 int turn;         // whose turn?
3
4 void init() {
5     flag[0] = flag[1] = 0;
6     turn = 0;
7 }
8 void lock(int self) {
9     flag[self] = 1;
10    turn = 1 - self;    // let other run
11    while ((flag[1-self] == 1) && (turn == 1-self))
12        ; // spin-wait
13 }
14 void unlock(int self) {
15     flag[self] = 0;
16 }
```

- Various issues → concurrency course

kalloc, good one

```
struct {  
    struct run *freelist;  
    int mutex;  
} kmem;  
  
void *kalloc(void)  
{  
    struct run *r;  
  
    hardLock(&kmem.mutex);  
    r = kmem.freelist;  
    if (r)  
        kmem.freelist = r->next;  
    hardUnlock(&kmem.mutex);  
    return (void*)r;  
}
```

Again. Is this good?

- It works.
- The interrupts are disabled for a SHORT period of time.
- The busy-wait (aka spin-lock) will run for a SHORT period of time.

So, it is reasonable. No better solution.

However. What happens if the critical section is one of the following:

- long
- required interrupt machinery (which anyways means long)

Easy now. Use `hardLock/hardUnlock` instead of `intDsbl/intEnbl`.

softLock

```
struct {  
    int haMutex;  
    int mutex;  
} softMutex;  
  
softLock(softMutex *sMutex) {  
    hardLock(&sMutex->haMutex);  
  
    while (s->mutex != 0) {  
        hardUnlock(&sMutex->haMutex);  
        BLOCK();  
        hardLock(&sMutex->haMutex);  
    }  
    sMutex->mutex = 1;  
  
    hardUnlock(&sMutex->haMutex);  
}
```