

I/O Devices - General (ch. 36+37)

Operating Systems

Based on: Three Easy Pieces by Arpaci-Dusseau

Moshe Sulamy

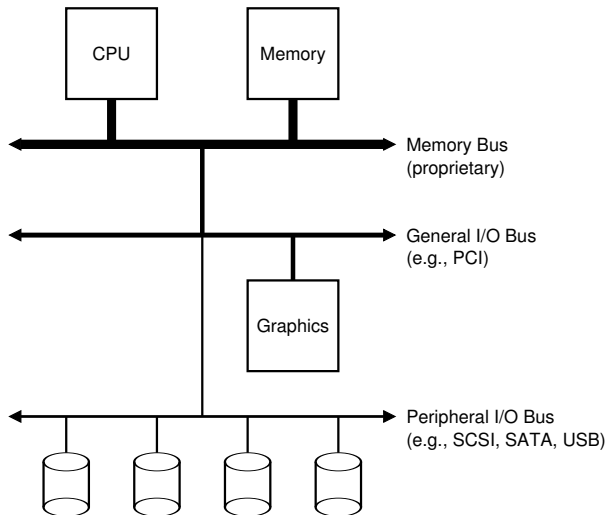
Tel-Aviv Academic College

I/O Devices

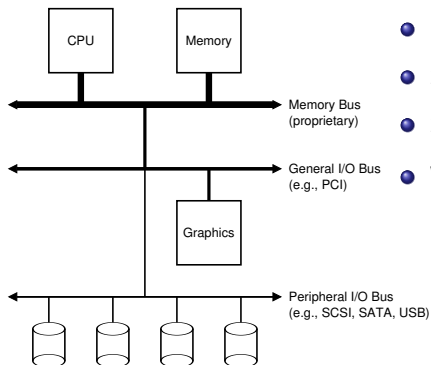
- New part: **persistence**
- But first: **input/output (I/O) devices**
 - Critical to computer systems

How should I/O be integrated into systems?

System Architecture




System Architecture



- CPU attached to memory via **memory bus**
- Some devices via **general I/O bus**
- Slow devices via **peripheral bus**
- Why hierarchical?
 - Physics and cost
 - Faster bus → shorter
 - Lower performance → further


Canonical Device

- A canonical device has two important components:
 - **Hardware interface:** allows the system to control its operation
 - **Internal structure:** implementation specific to the device



Canonical Protocol

- Device interface comprised of three registers
 - **Status**: current status of device
 - **Command**: tell device to perform a task
 - **Data**: pass data to device or get data from it



- Control device behavior by reading and writing these registers

Canonical Protocol

- Typical interaction of OS with the device:

```
1 while (STATUS == BUSY)
2     ; // wait until device is not busy
3 write data to DATA register
4 write commands to COMMAND register
5     // starts the device and executes the command
6 while (STATUS == BUSY)
7     ; // wait until device is done with your request
```

Canonical Protocol

① Polling

- Repeatedly reading status register

② OS sends some data

- Multiple writes may be needed
- CPU involved with data movement: **programmed I/O (PIO)**

③ Write command


- Lets device know that data is present

④ Polling

- OS waits for device to finish

Canonical Protocol


- Polling is inefficient
 - Wastes CPU time waiting for device
 - Switch to another process: better utilize CPU



How can the OS check device status without polling?

Interrupts


- Instead of pulling, issue a request
 - Put calling process to sleep
 - Context switch to another task
- Device finished: hardware interrupt
 - CPU jumps into OS **interrupt handler**
 - (Also: **interrupt service routing (ISR)**)
 - Handler will finish the request and wake waiting process




Interrupts

- Not always the best solution
 - Device performs very quickly
 - Interrupts will slow down the system
 - Switching back and forth is expensive
- **Hybrid (two-phased approach)**
 - Poll for a little while
 - If not finished, use interrupts

- Programmed I/O:
 - CPU transfers a large chunk of data to a device
 - CPU overburdened with a trivial task



- Solution: **Direct Memory Access (DMA)**
 - DMA controller (a device) handles copying of data
 - OS programs DMA:
 - Where the data lives in memory
 - How much data to copy
 - Which device to send to/read from



Device Interaction

- How does the OS/CPU communicate with devices?
- **I/O instructions**
 - OS sends data to specific device registers
 - For example, `in` and `out` **privileged** instructions on x86
- **Memory-mapped I/O**
 - Device registers available as if they were memory locations
 - OS issues `load` or `store` to address
 - Hardware routes to device instead of main memory

Device Driver

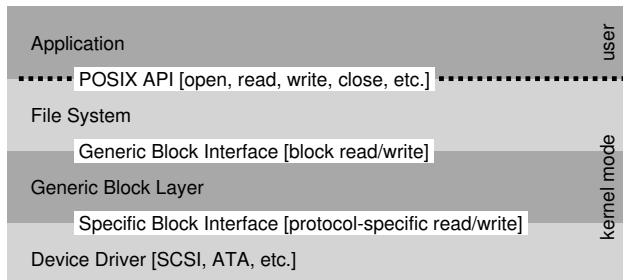
- Devices have specific interfaces
 - Keep OS as general as possible
 - e.g., build file system that works on SCSI disks, IDE disks, USB drives, etc.
- Solution?

Device Driver

- Devices have specific interfaces
 - Keep OS as general as possible
 - e.g., build file system that works on SCSI disks, IDE disks, USB drives, etc.
- Solution? **abstraction**
 - Software that knows device specifics: **device driver**
 - Over 70% of OS code in Linux
 - Primary contributor to **kernel crashes**

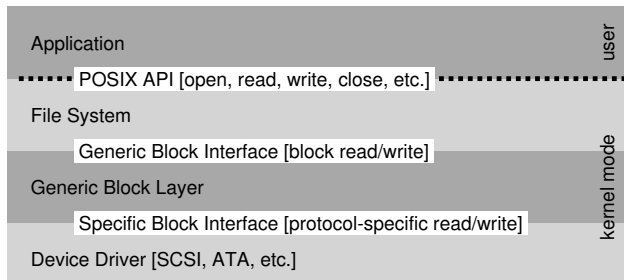
Device Driver

- The Linux file system software stack:



Device Driver

- The Linux file system software stack:



- Also available: **raw interface**
 - Enables special applications to directly read and write blocks
 - e.g., file-system checker, disk defragmentation tool

Summary (I/O Devices)

- Canonical device: registers, HW interface, internal structure
- Canonical protocol: polling, data, command, polling
- Interrupts: instead of polling, issue a request
 - Device finished: hardware interrupt
 - **Hybrid** approach: poll for a little while, then use interrupts
- **DMA** controller: handles copying of data
- Device interaction: **I/O instructions** or **memory-mapped I/O**
- **Device driver**: software abstraction that knows device specifics